



Azure Cosmos DB

Technical Overview

Chris Joakim, Specialist / Global Black Belt, Charlotte

Hailey Wu, Cloud Solution Architect, Toronto

Xavier Guérette, Cloud Solution Architect, Montreal



AGENDA

Overview

- CosmosDB: What it is, and why use it?
- Resource Model, Provisioning, Automation

The Basics

- System Topology & Partitioning
- Request Units (RU) and Scaling
- Global Distribution & Consistency Models
- HA/DR, Backups
- Logging / Azure Monitor / Kusto

Main Features

- Change Feed
- TTL

Main Features (continued)

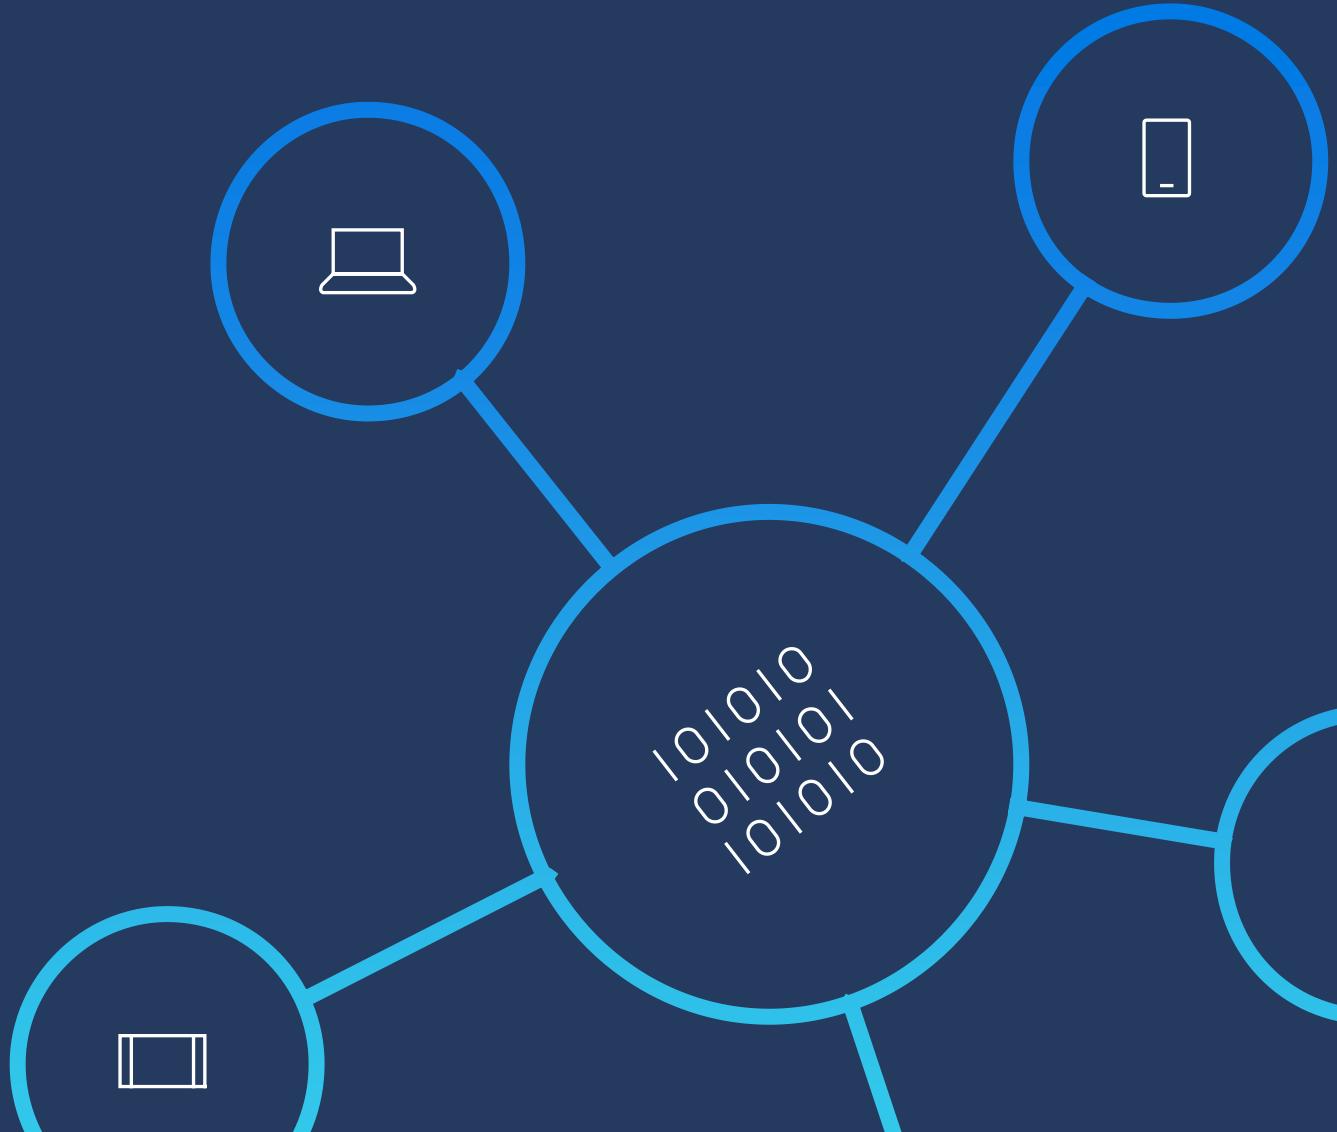
- HTAP with Synapse Link
- Integrated Cache
- Integration Examples:
 - Streaming – Stream Analytics, Functions
 - Search with Azure Cognitive Search

Questions, Discussion, Next Steps

Design and Development (Next Session)

- Design – patterns and anti-patterns
- Java – Spring, Spring Data, SDK
- Links to Documentation and Examples

OVERVIEW



Cosmos DB

- A Family of NoSQL Databases
- PaaS / Managed
- Scalable
- Distributed
- HA/DR, SLAs
- Secure
- Integrations

APIs

- **Core (SQL) API** (msft)
- Table (msft)
- MongoDB (oss)
- Cassandra (oss)
- Gremlin (oss)



Azure Cosmos DB

SQL

SQL

{LEAF}

API for MongoDB



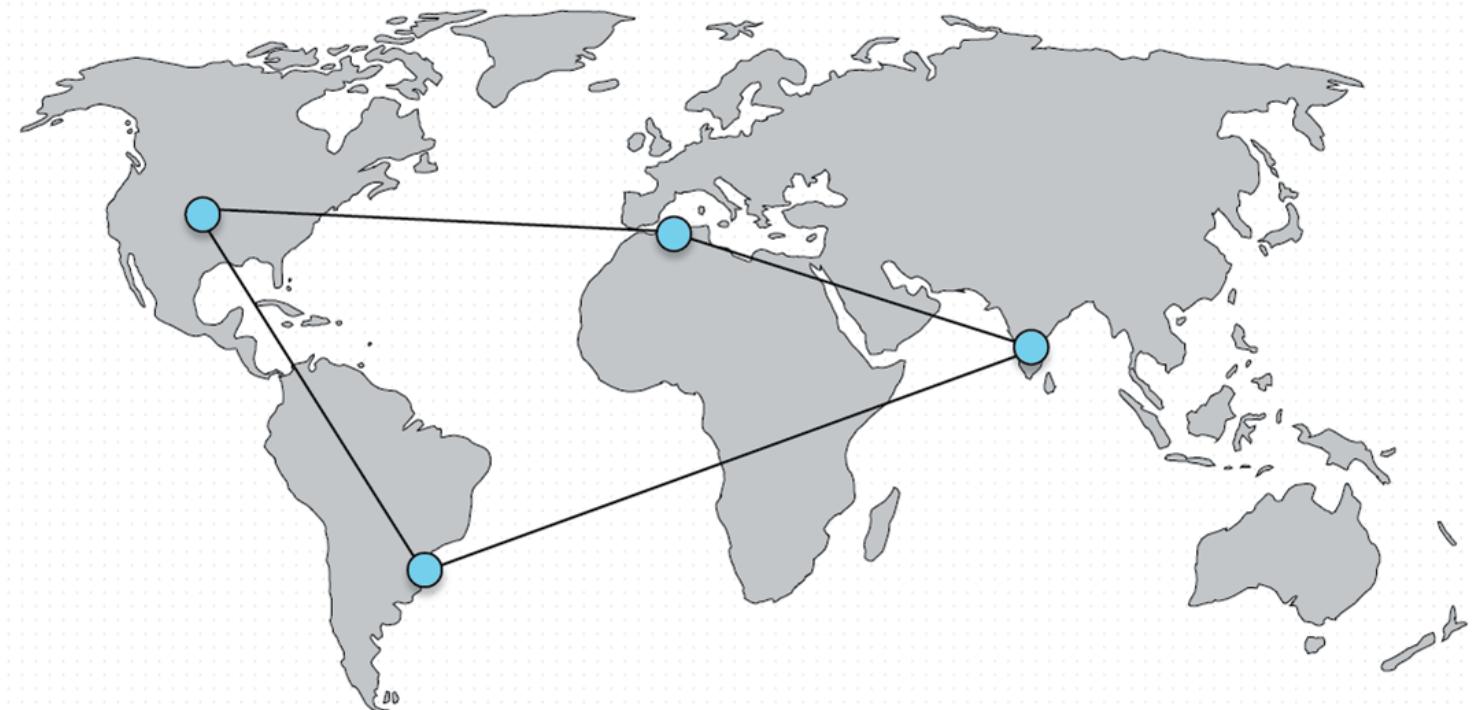
Gremlin

eye

Cassandra



Table



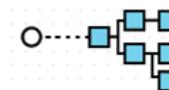
Key-Value



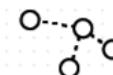
Column-Family



Documents



Graph



Guaranteed speed at any scale

Simplified application development

Mission-critical ready

Fully managed and cost effective

TOP 10 REASONS WHY CUSTOMERS USE AZURE COSMOS DB



The 1st and only database with **global distribution turnkey capability**



Very small to massively scalable in the same account



Guaranteed single digit millisecond latency at 99th percentile



Excellent integration with other Azure PaaS services



Boasts 5 well-defined consistency models to pick the right consistency/latency/throughput tradeoff



High SLA, suitable for mission critical applications



High flexibility to optimize for speed and cost. RU cost model.



Tackles **big data** workloads with **high availability and reliability**



Enterprise-grade security. In flight and on-disk encryption, BYOK, RBAC, Private Networking, etc.



Perfect for **event-driven** architectures as well as **analytics**. Ease of use.

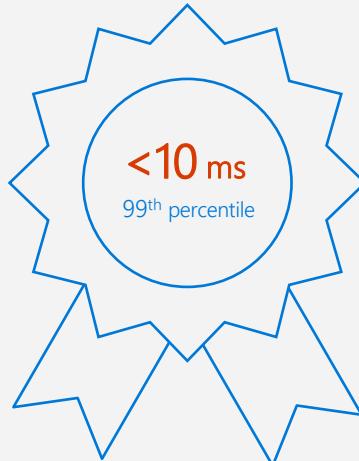
COMPREHENSIVE SLAS

COVERS AVAILABILITY AND LATENCY

RUN YOUR APP ON WORLD-CLASS INFRASTRUCTURE

Azure Cosmos DB is the only service with **financially-backed SLAs for single-digit millisecond read and write latency** at the 99th percentile, **99.999% high availability** and guaranteed throughput and consistency

Latency



High Availability



Throughput



Consistency



GROWTH MINDSET; ALWAYS BE LEARNING PLEASE EMBRACE LEARNING COSMOSDB TO LEVERAGE IT BEST

Similar, but different:

- C# != Java
- CosmosDB != DB2
- CosmosDB != Atlas

Each programming language and DB has its' own Idioms/Patterns.

Please take the time to learn CosmosDB; it is different.

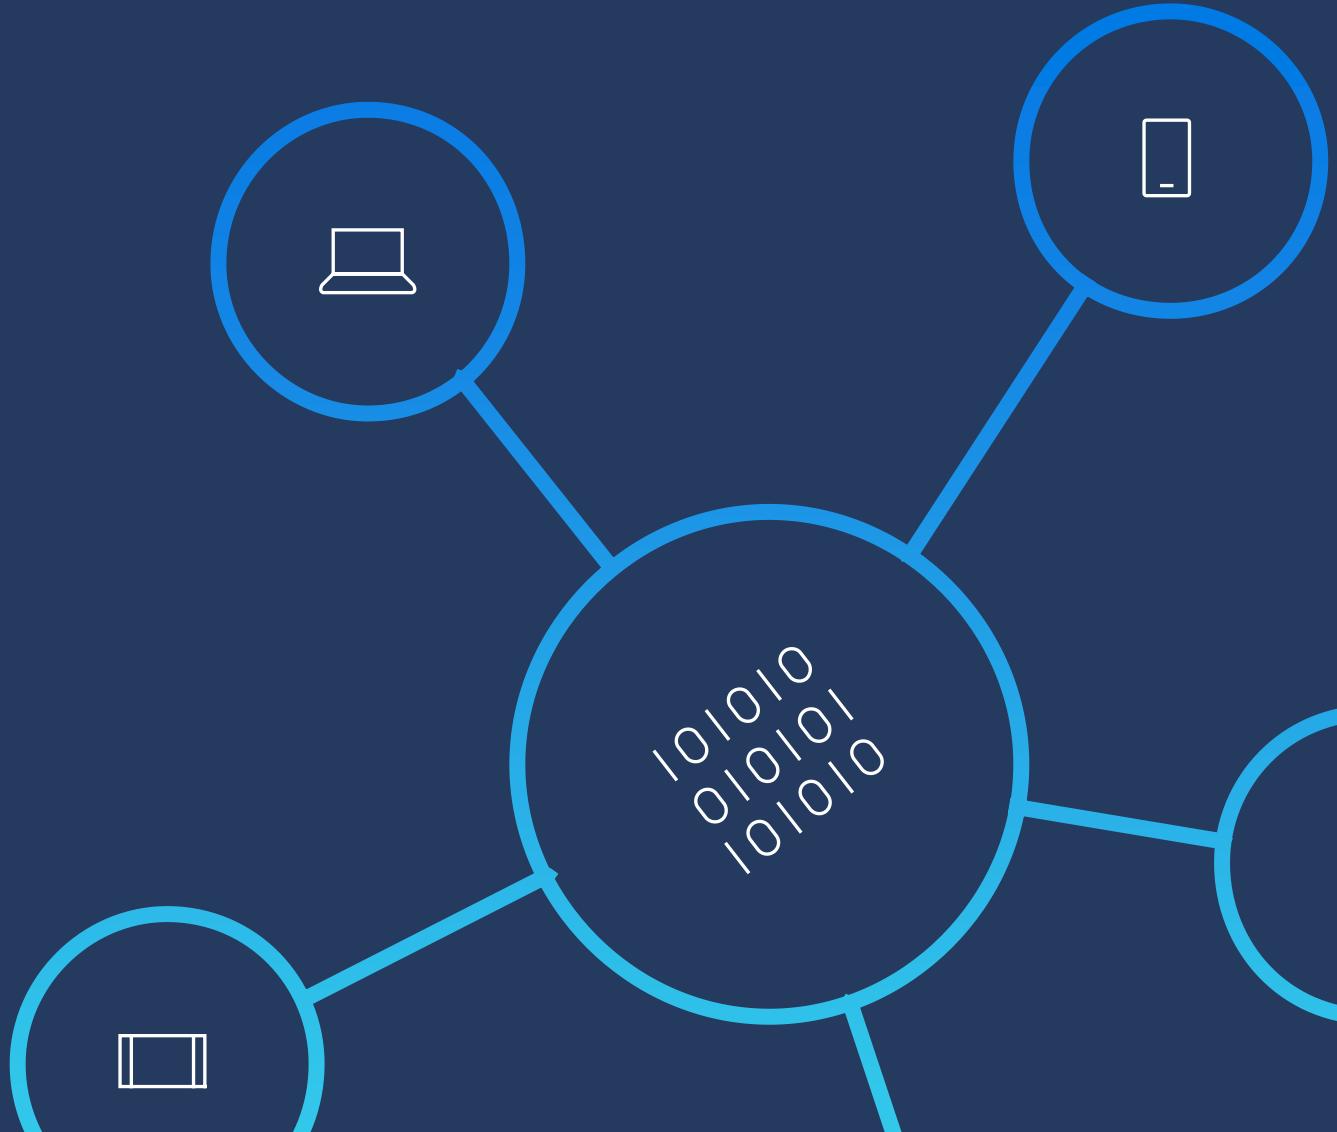
Growth Mindset

"We need to be always learning and insatiably curious. We need to be willing to lean in to uncertainty, take risks and move quickly when we make mistakes, recognizing failure happens along the way to mastery."

- Satya Nadella

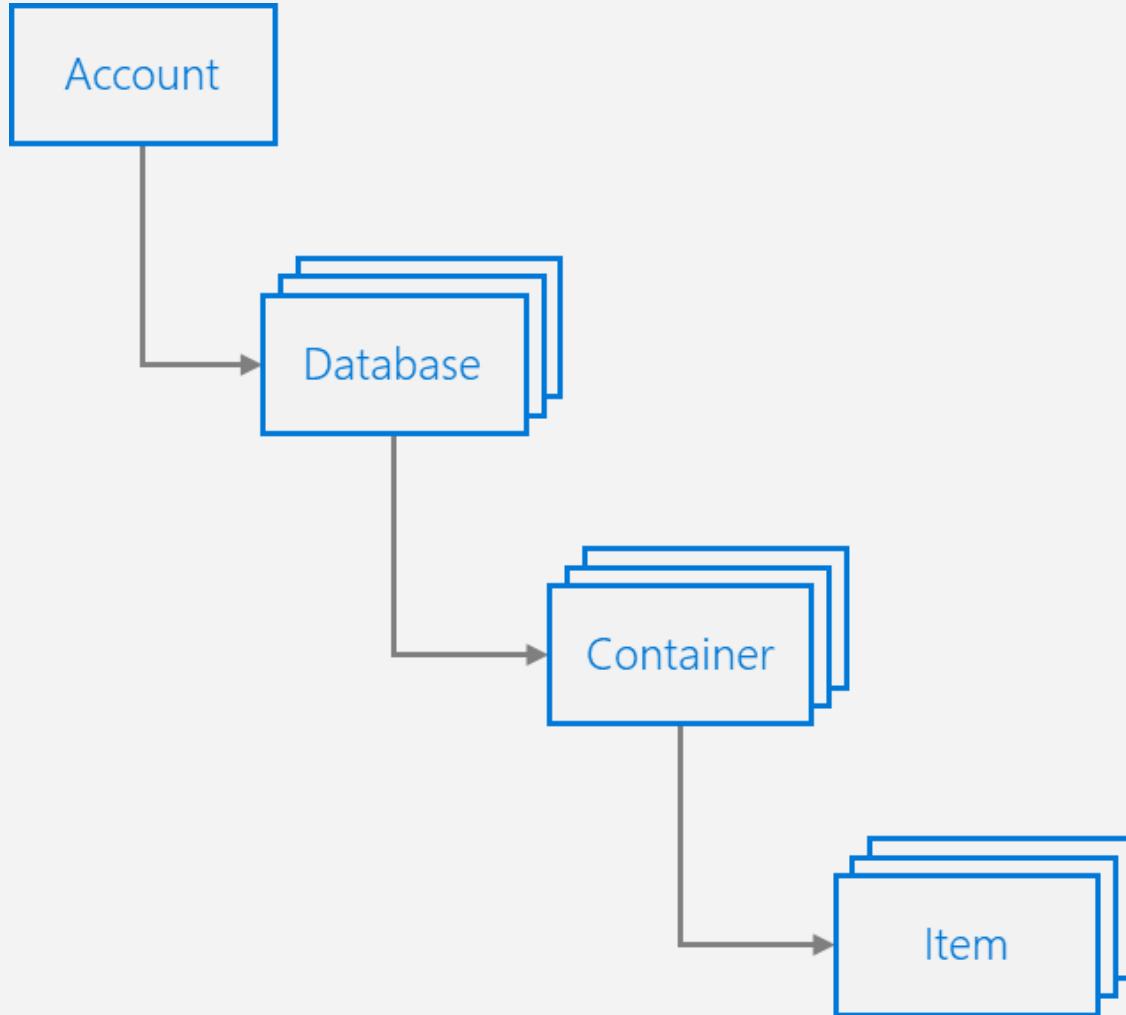


The Basics



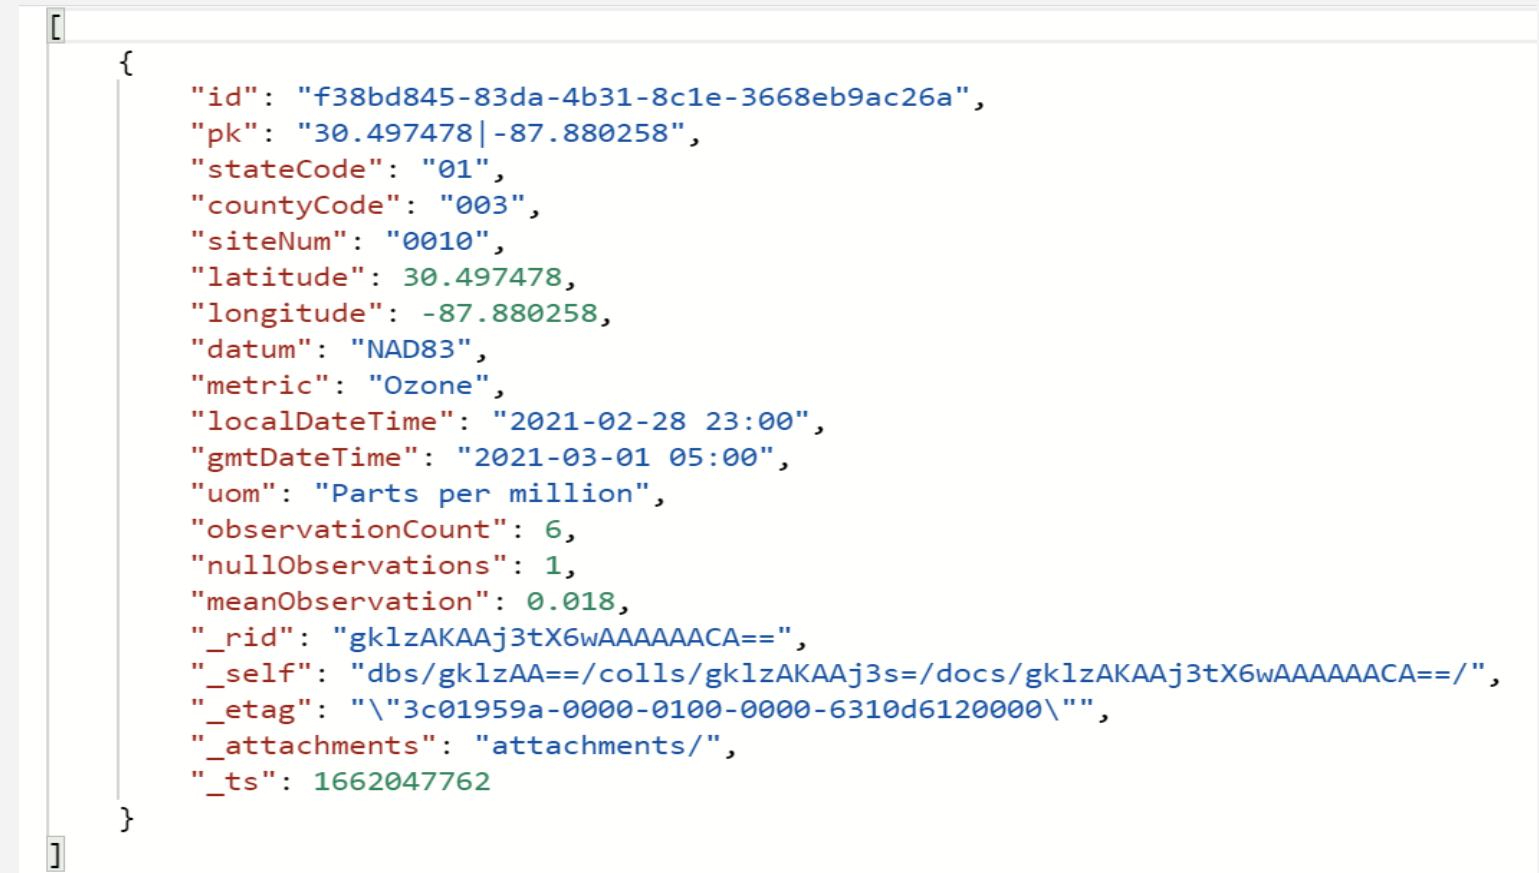
THE STRUCTURE OF A COSMOSDB ACCOUNT

"CONTAINER" == "COLLECTION"



IS THE COSMOSDB SQL API RELATIONAL? NO!

- The data is JSON
- It's Schemaless
 - other than the pk
- Nesting
- Query it with SQL
- Indexing – default, custom, paths, composite



A screenshot of a code editor showing a single JSON object. The object has an empty array at the top level, followed by a brace '{' and a series of key-value pairs. The keys are color-coded: 'id' (blue), 'pk' (red), 'stateCode' (blue), 'countyCode' (blue), 'siteNum' (blue), 'latitude' (green), 'longitude' (green), 'datum' (blue), 'metric' (blue), 'localDateTime' (blue), 'gmtDateTime' (blue), 'uom' (blue), 'observationCount' (blue), 'nullObservations' (blue), 'meanObservation' (green), '_rid' (blue), '_self' (blue), '_etag' (blue), '_attachments' (blue), and '_ts' (blue). The values are mostly strings, with some numbers like 0.018 and 1662047762.

```
[{"id": "f38bd845-83da-4b31-8c1e-3668eb9ac26a", "pk": "30.497478|-87.880258", "stateCode": "01", "countyCode": "003", "siteNum": "0010", "latitude": 30.497478, "longitude": -87.880258, "datum": "NAD83", "metric": "Ozone", "localDateTime": "2021-02-28 23:00", "gmtDateTime": "2021-03-01 05:00", "uom": "Parts per million", "observationCount": 6, "nullObservations": 1, "meanObservation": 0.018, "_rid": "gklzAKAAj3tX6wAAAAAACAA==", "_self": "dbs/gklzAA==/colls/gklzAKAAj3s=/docs/gklzAKAAj3tX6wAAAAAACAA==/", "_etag": "\"3c01959a-0000-0100-0000-6310d6120000\"", "_attachments": "attachments/", "_ts": 1662047762}]
```

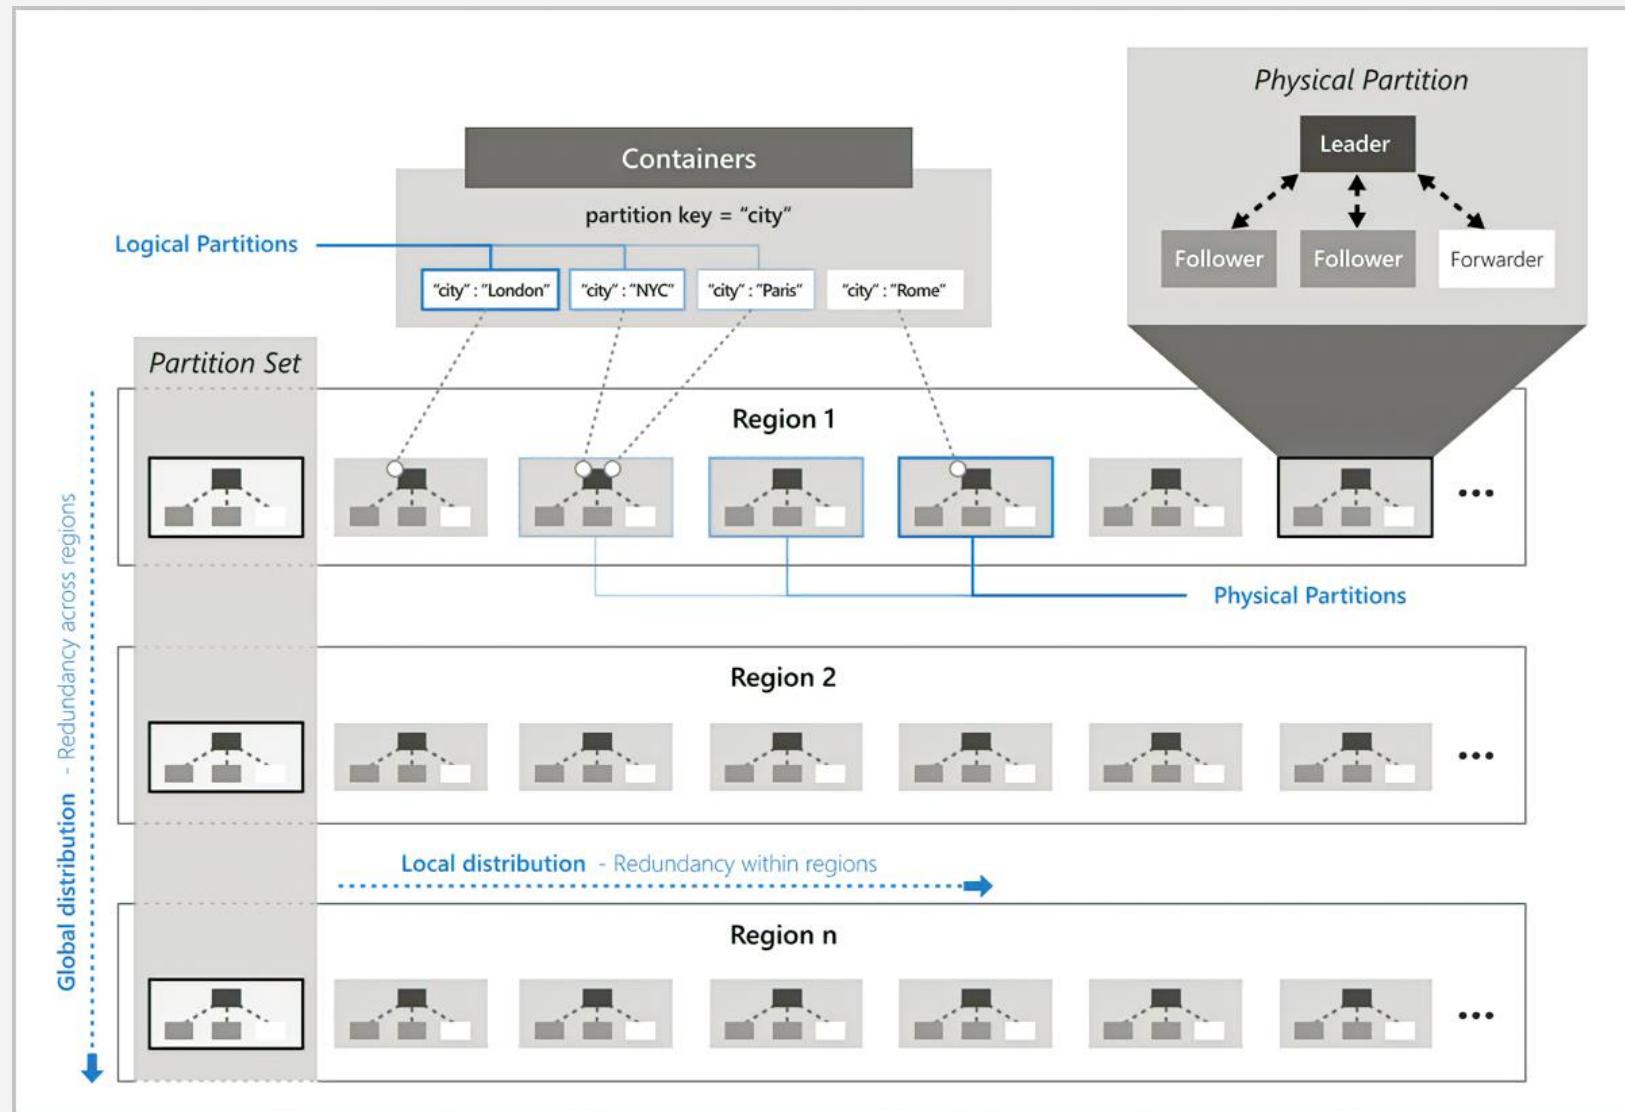
```
SELECT * FROM c WHERE c.id = 'f38bd845-83da-4b31-8c1e-3668eb9ac26a' AND c.pk = '30.497478|-87.880258'
```

REQUEST UNITS – THE UNIT OF SCALE AND PRICE

- Request Units (RU) are a scalable unit of Memory, CPU, and IOPS
- RUs are the primary cost component of CosmosDB
 - storage, network egress, backups, etc. are other cost components
- The cost to read a **1.0 KB document** by Id and Partition Key is **1.0 RU**
- Approx 5x for writes
- RUs can either be **Manually or Auto-Scaled**, at the container or DB level
- Think of Request Units as a configured **“Per-Second Budget”** of Throughput
- Think of an Uber ride vs buying a car; you pay per use
- **RU costs are provided by our SDKs, Azure Portal, Azure Monitor**
- “Burst” Capacity functionality is currently in preview mode
- **The RU cost model has design implications for your application**
 - The size of your documents matters, as well as the updates
- <https://docs.microsoft.com/en-us/azure/cosmos-db/request-units>

HORIZONTAL SCALING, PHYSICAL AND LOGICAL PARTITIONING

EACH CONTAINER HAS A FIXED PARTITION KEY ATTRIBUTE



REPLICATE TO ANY AZURE REGION(S)

USE A SINGLE REGION, OR MULTIPLE WRITE REGIONS (W/ FAILOVER CONFIG)

Replicate data globally

... [Save](#) [Discard](#) [Manual failover](#) [Service-Managed Failover](#)

⚠️ When you add a region to your account, you will be billed for the additional RU/s and storage copied to the region. Click here to learn more. →

Click on a location to add or remove regions from your Azure Cosmos DB account.

* Each region is billable based on the throughput and storage for the account. [Learn more](#)

Configure regions

Multi-region writes [i](#)

[Disable](#) [Enable](#)

Configure the regions for reads, writes and availability zone
[+ Add region](#)

Write region

East US

Read regions

North Europe

FIVE CONSISTENCY LEVELS FOR MULTIPLE WRITE REGIONS

Default consistency ...



Save



Discard

STRONG

BOUNDED STALENESS

SESSION

CONSISTENT PREFIX

EVENTUAL

HIGH AVAILABILITY, DISASTER RECOVERY, BACKUPS

- Periodic or Continuous Backups
- Restores to another account
- Periodic Backups
 - You specify interval and retention
 - Support Ticket to execute a restore
- Continuous Backups
 - Granular in time
 - You execute the restore (PITR)
- Also, Synapse Link (discussed later)

Backup & Restore ...

Your account is on periodic backup mode. You can now change to continuous mode for a better backup and restore experience. [Change to continuous mode](#)

Backup policy mode (change)
Periodic

Backup Interval
How often would you like your backups to be performed?
240 Minute(s) 60-1440

Backup Retention
How long would you like your backups to be saved?
8 Hours(s) 8-720

Copies of data retained 2

Backup storage redundancy * ⓘ
 Geo-redundant backup storage
 Zone-redundant backup storage
 Locally-redundant backup storage

LOGGING, MONITORING, ALERTS

SLIDE 1 – EXAMPLE CLASSIC (ORIGINAL) INHERENT METRICS

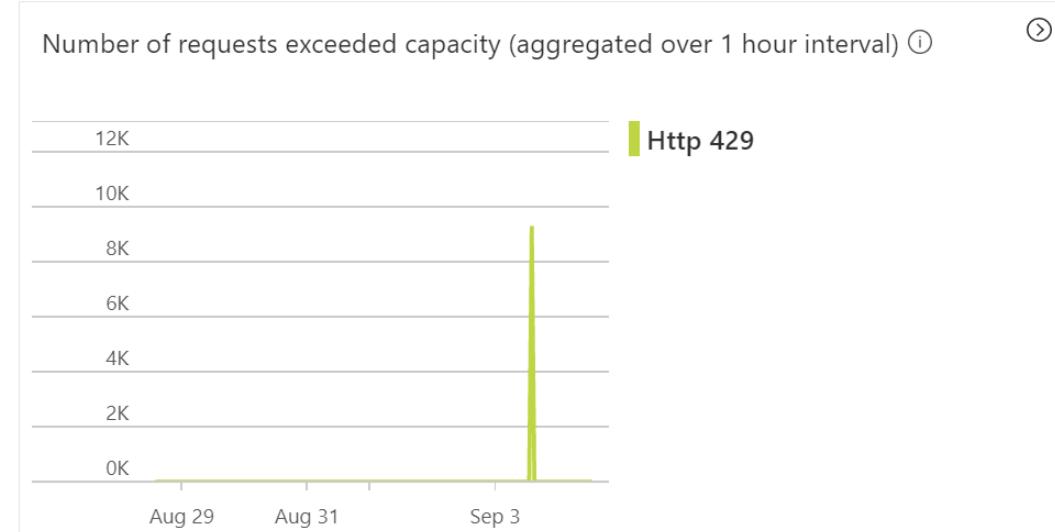
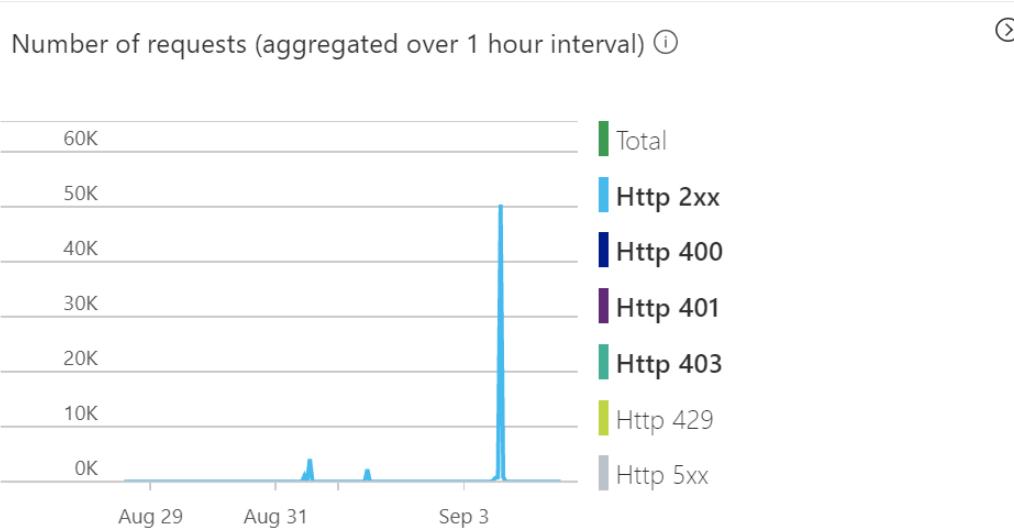
Metrics (Classic) ...

[Download as csv](#) [Refresh](#) [Feedback](#)

Overview **Throughput** Storage Availability Latency Consistency System

Database(s) Container(s) Region(s) Show data for last

Collection telemetry has provisioned throughput at the database level.



LOGGING, MONITORING, ALERTS

SLIDE 2 – AZURE MONITOR - WHAT GETS LOGGED, AND TO WHERE

Dashboard > cjoakimcosmosql | Diagnostic settings > Diagnostic settings >

Diagnostic setting

Save Discard Delete Feedback

A diagnostic setting specifies a list of categories of platform logs and/or metrics that you want to collect from a resource, and one or more destinations that you would stream them to. Normal usage charges for the destination will occur. [Learn more about the different log categories and contents of those logs](#)

Diagnostic setting name cosmssql-la

Logs

Categories

- DataPlaneRequests
- MongoRequests
- QueryRuntimeStatistics
- PartitionKeyStatistics
- PartitionKeyRUConsumption
- ControlPlaneRequests
- CassandraRequests
- GremlinRequests
- TableApiRequests

Metrics

- Requests

Destination details

Send to Log Analytics workspace

Subscription Microsoft Azure Internal Consumption

Log Analytics workspace cjoakimla (eastus)

Destination table ①
Azure diagnostics Resource specific

Archive to a storage account

Stream to an event hub

Send to partner solution

LOGGING, MONITORING, ALERTS

SLIDE 3 – AZURE MONITOR QUERIES WITH THE KUSTO SYNTAX

- How to get partition key statistics to evaluate skew across top 3 partitions for a database account:

Kusto

 Copy

```
CDBPartitionKeyStatistics  
| project RegionName, DatabaseName, CollectionName, PartitionKey, SizeKb
```

- How to get the request charges for expensive queries?

Kusto

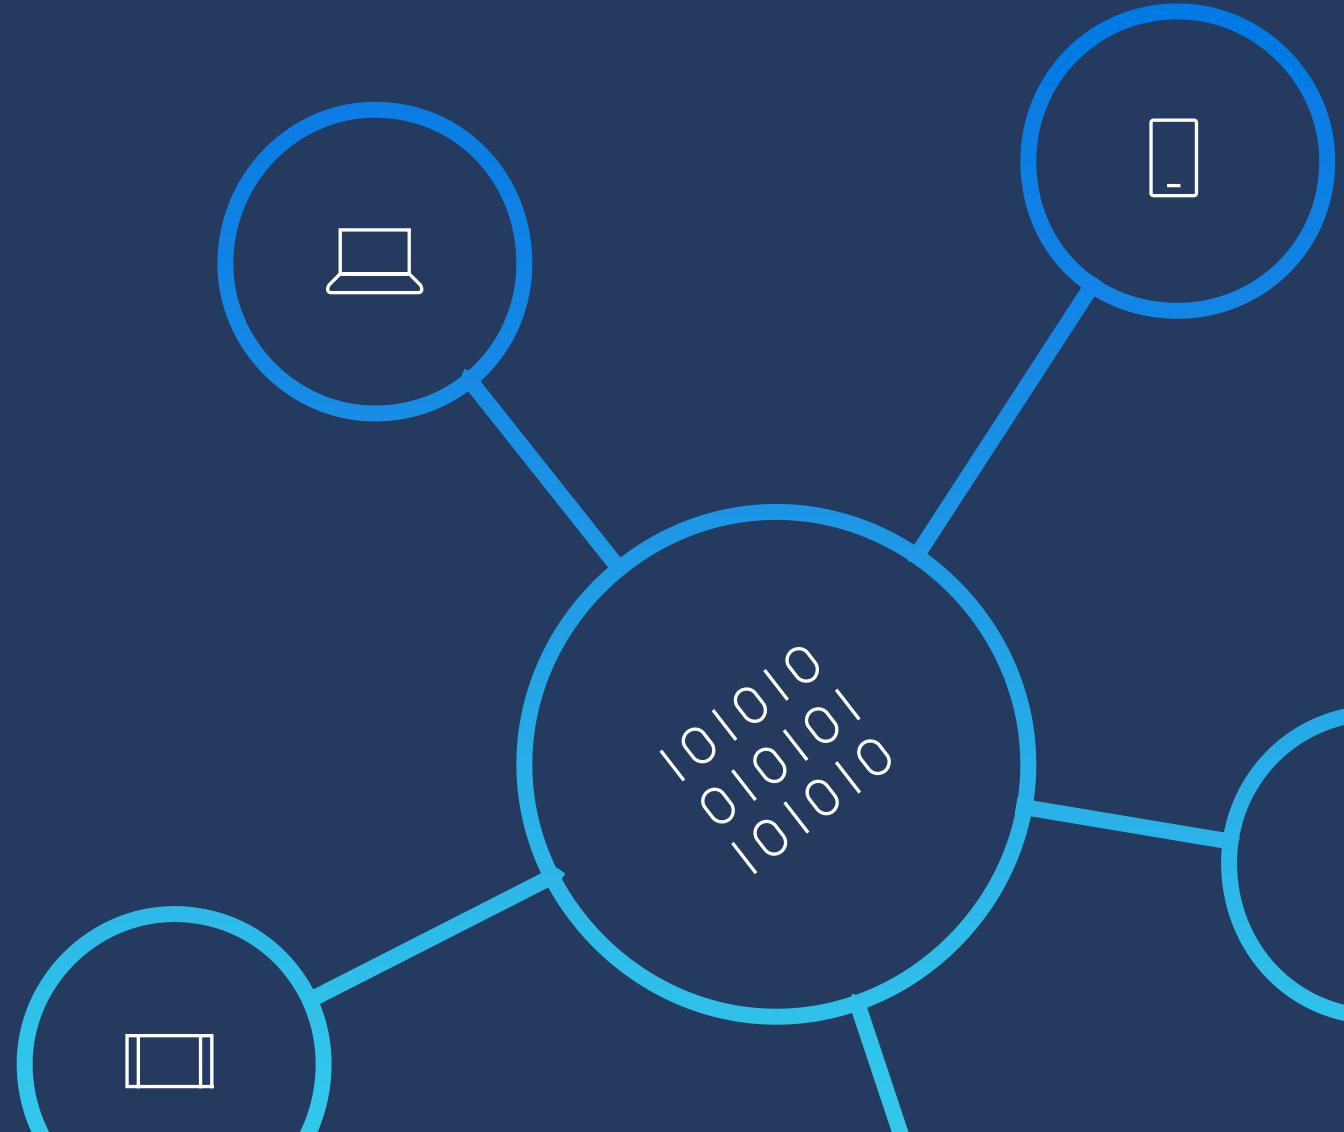
 Copy

```
CDBDataPlaneRequests  
| where todouble(RequestCharge) > 10.0  
| project ActivityId, RequestCharge  
| join kind= inner (  
CDBQueryRuntimeStatistics  
| project ActivityId, QueryText  
) on $left.ActivityId == $right.ActivityId  
| order by RequestCharge desc  
| limit 100
```

<https://docs.microsoft.com/en-us/azure/cosmos-db/cosmosdb-monitor-logs-basic-queries>

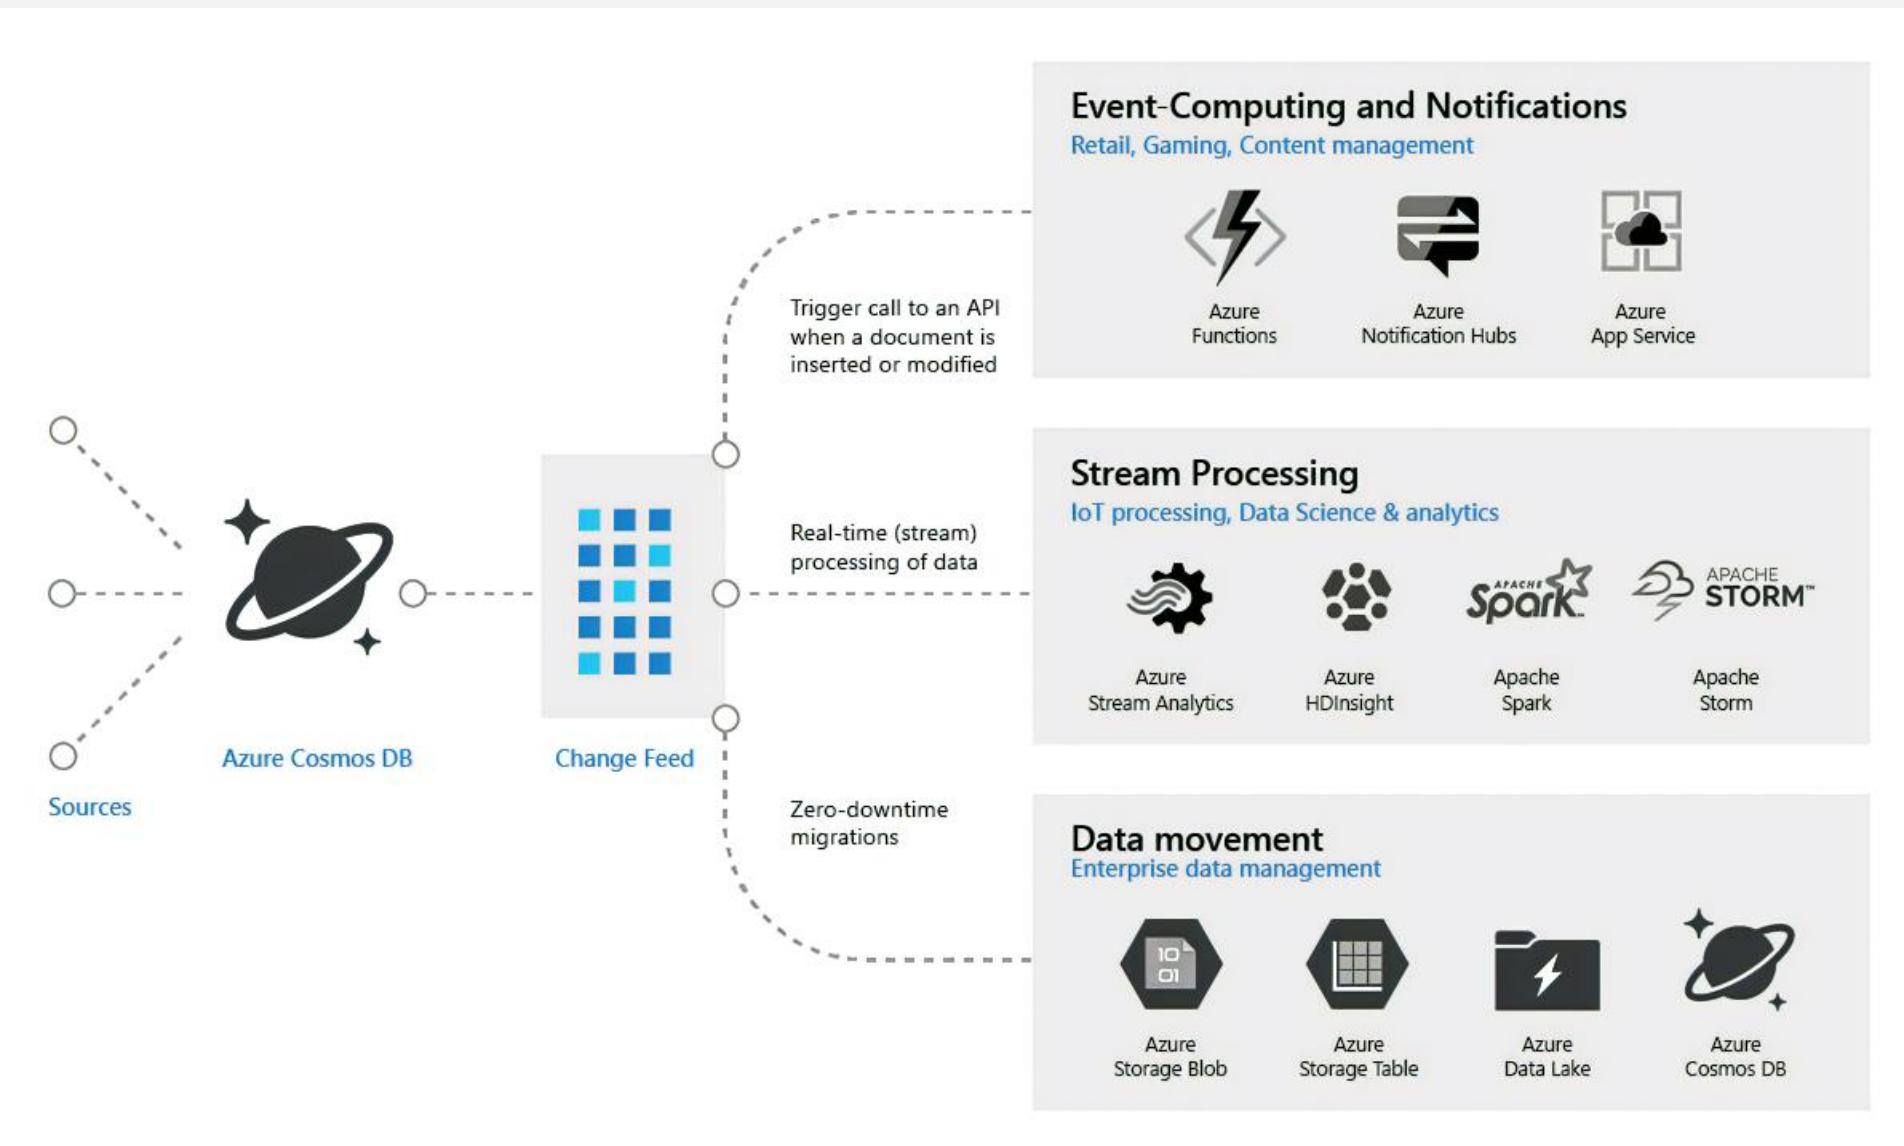
<https://docs.microsoft.com/en-us/azure/cosmos-db/cosmos-db-advanced-queries?tabs=resource-specific>

Main Features



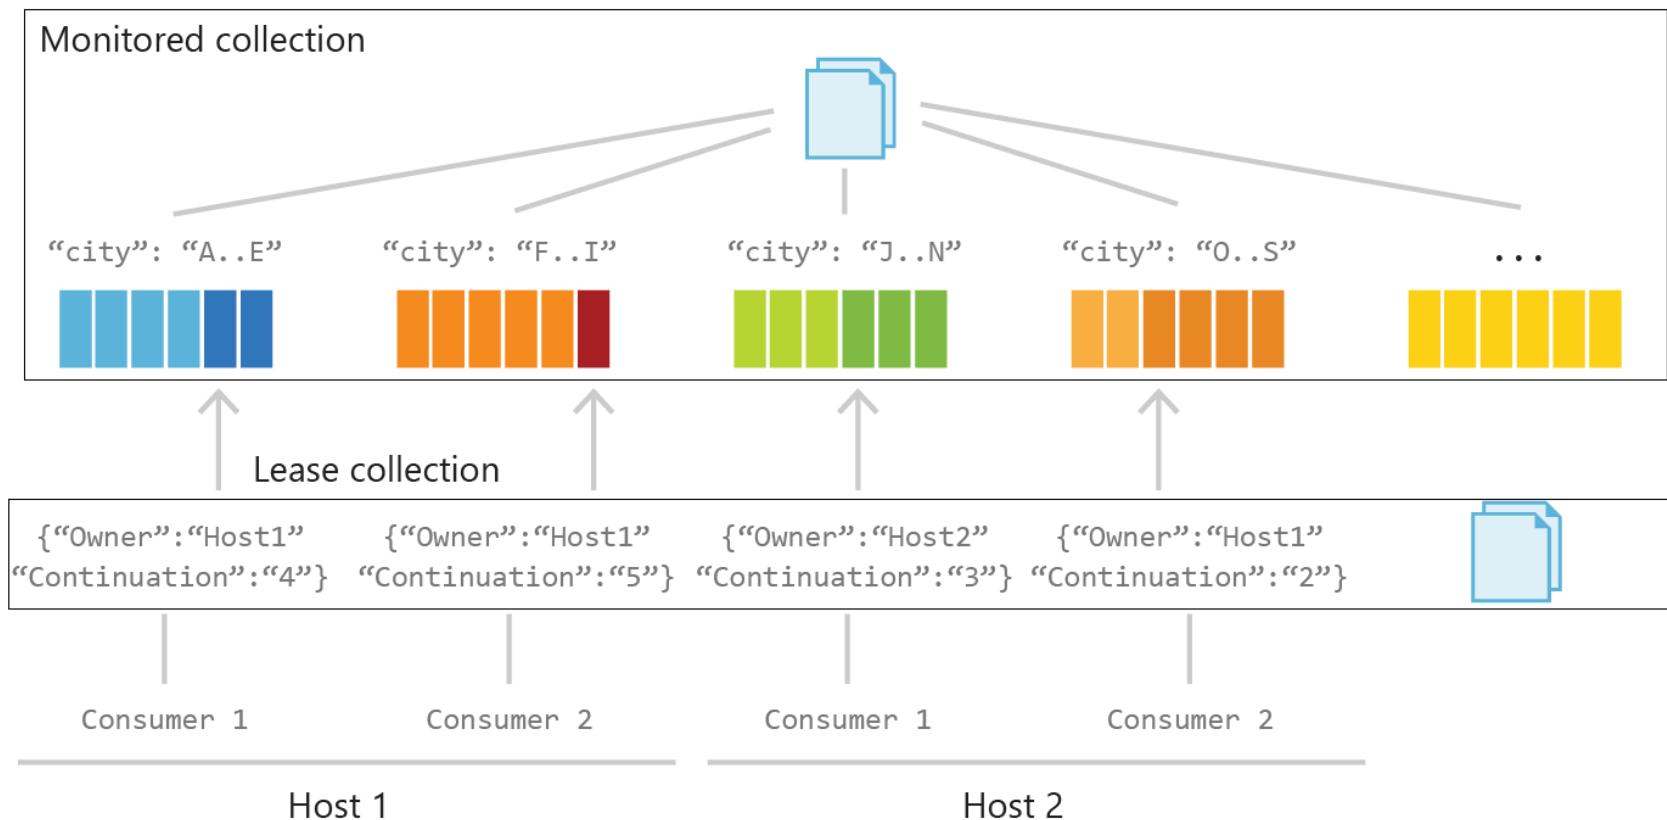
COSMOSDB CHANGE FEED

CONSUME WITH AZURE FUNCTIONS OR A SDK (JAVA OR C#)



CHANGE FEED PROCESSOR PARTITIONS, CONSUMERS, LEASES

A “Leases” container maintains the state of a change-feed consumer



TIME-TO-LIVE (TTL) AUTOMATIC DELETION WITH ZERO RU COSTS, BASED ON _TS

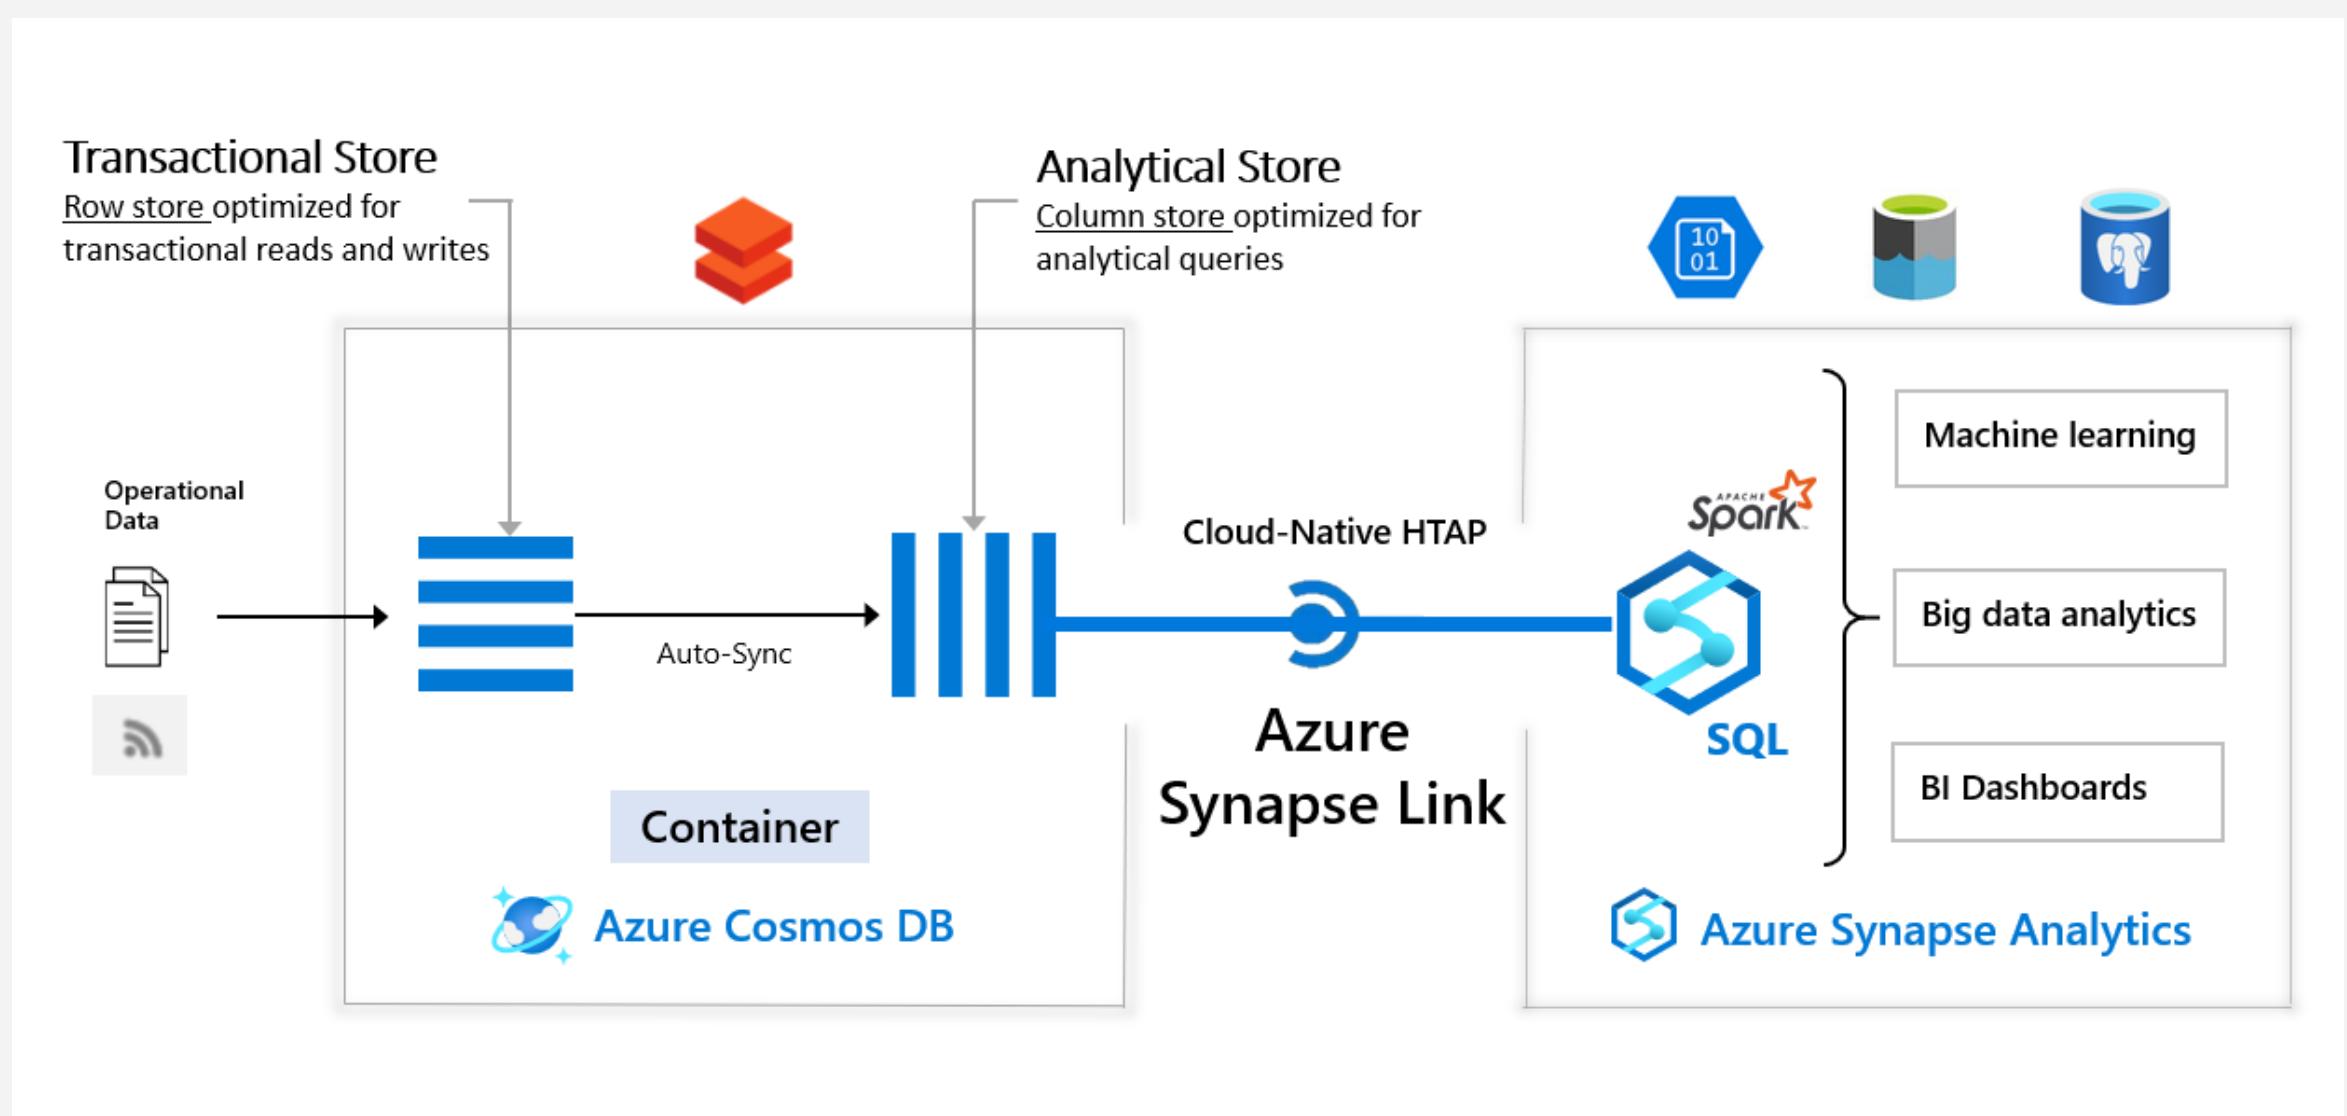
Data Explorer

The screenshot shows the Azure Data Explorer interface. The left sidebar lists datasets: DATA, dev, telemetry, testing, users, graph, and graph. The telemetry dataset is selected. The main area shows the 'Settings' tab for the telemetry dataset. Under 'Time to Live', the 'On' radio button is selected, and the value '2592000 second(s)' is entered. Other options include 'Off' and 'On (no default)'. Below this is the 'Geospatial Configuration' section, where 'Geography' is selected. The top navigation bar includes icons for SQL API, JSON API, Log Analytics, and Kusto Query Editor, along with Save and Discard buttons.

```
1 {  
2   "id": "726a44ab-856f-4027-b429-312c30250a79",  
3   "pk": "33.904039|-86.053867",  
4   "stateCode": "01",  
5   "countyCode": "055",  
6   "siteNum": "0011",  
7   "latitude": 33.904039,  
8   "longitude": -86.053867,  
9   "datum": "NAD83",  
10  "metric": "Ozone",  
11  "localDateTime": "2021-02-28 23:00",  
12  "gmtDateTime": "2021-03-01 05:00",  
13  "uom": "Parts per million",  
14  "observationCount": 6,  
15  "nullObservations": 1,  
16  "meanObservation": 0.029,  
17  "_rid": "gklzAKAAj3tvNwAAAAAAA==",  
18  "_self": "dbs/gklzAA==/colls/gklzAKAAj3s=/docs/gklzAKAAj3tvNwAAAAAAA==/",  
19  "_etag": "\"9103d226-0000-0100-0000-6313a0960000\"",  
20  "_attachments": "attachments/",  
21  "_ts": 1662230678  
22 }
```

HTAP WITH SYNAPSE LINK

HYBRID TRANSACTIONAL ANALYTICAL PROCESSING (OLTP & OLAP)



HTAP WITH SYNAPSE LINK

READING THE ANALYTIC DATASTORE WITH SPARK IN AZURE SYNAPSE

```
1 # Load the SynapseLink Sales Data into a Dataframe.
2 # Select just the "sale" document types from the sales container,
3 # which have a minimum _ts (timestamp) value
4
5 from pyspark.sql.functions import col
6
7 # initialize variables
8 begin_timestamp = 0
9 end_timestamp   = 1699999999
10
11 # read just the doctype "sales", not "line_item"
12 # "cosmos.oltp" = CosmosDB live database
13 # "cosmos.olap" = Synapse Link Analytic Datastore
14
15 df_sales = spark.read\
16     .format("cosmos.olap")\
17     .option("spark.synapse.linkedService", "CosmosSqlDemoDB")\
18     .option("spark.cosmos.container", "sales")\
19     .load().filter(col("doctype") == "sale")\
20     .filter(col("_ts") > begin_timestamp)\
21     .filter(col("_ts") < end_timestamp)
22
23 display(df_sales.limit(3))
```

INTEGRATED CACHE

GATEWAY (HTTPS) OR DIRECT (TCP)

Connection Modes:

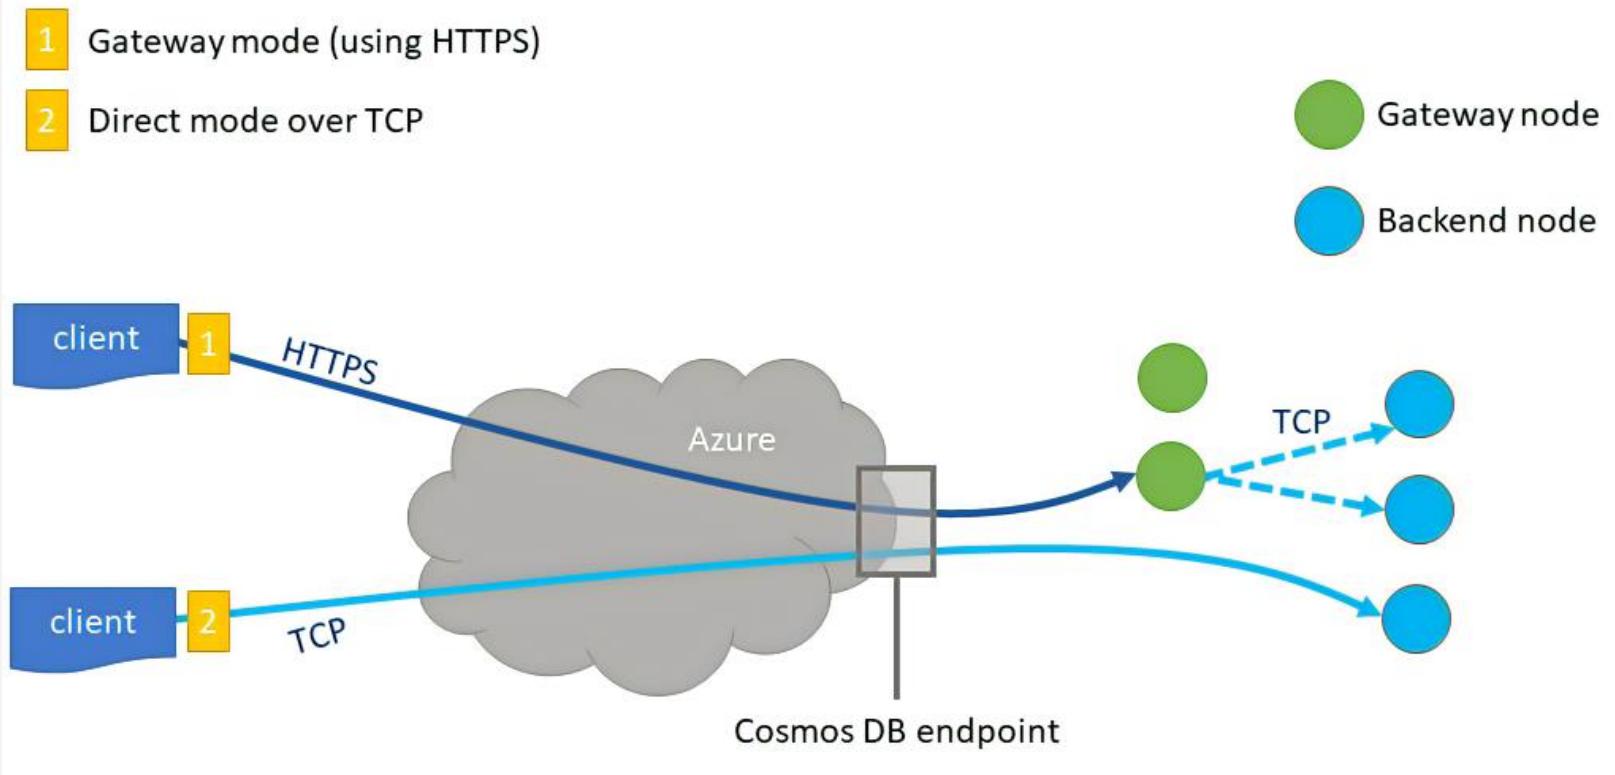
- HTTPS (Gateway)
- TCP (Direct)

Gateways:

- Standard (default)
- **Dedicated (cache)**

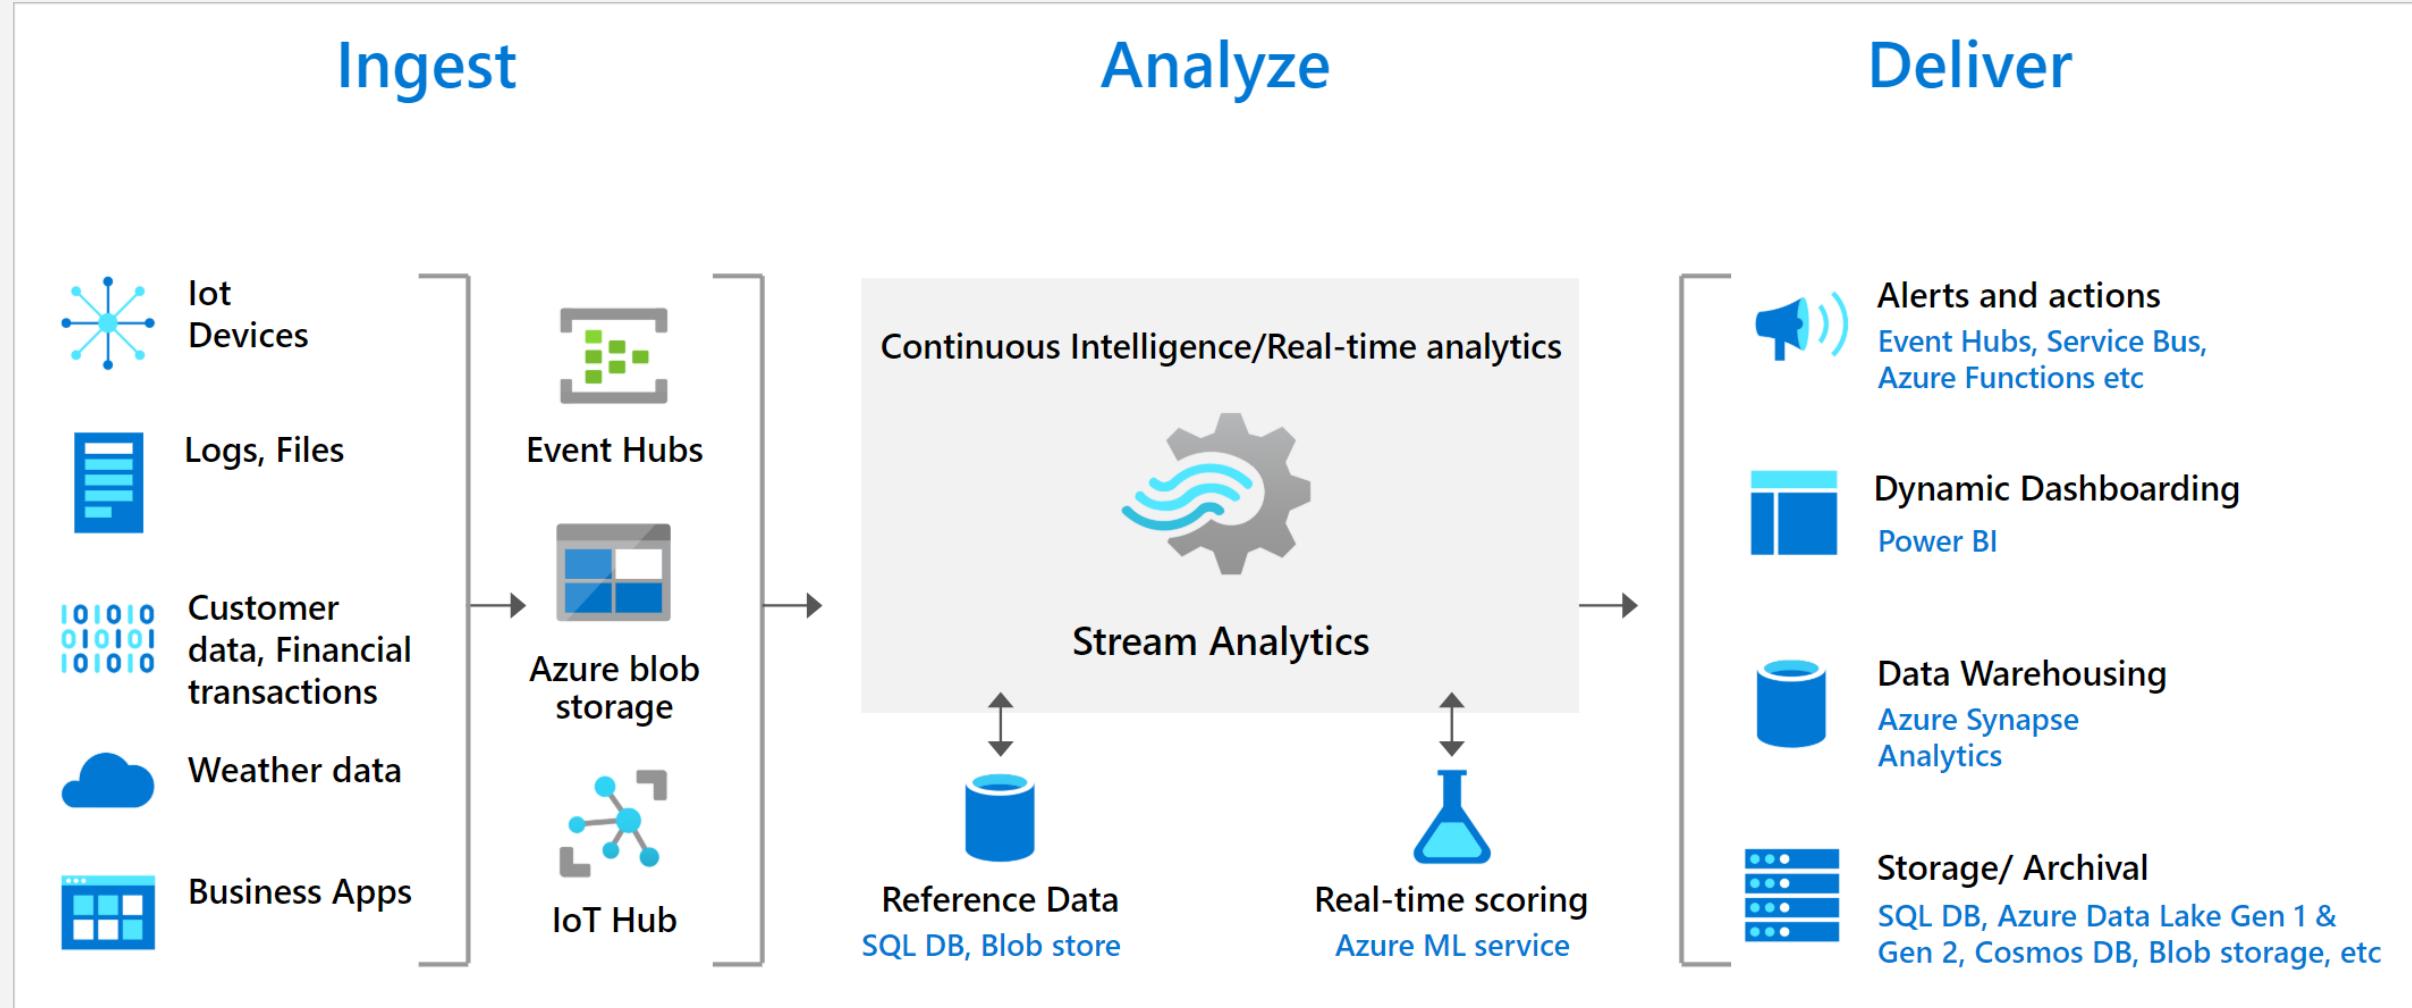
Caches:

- Sizes – S, M, L
- Types – Query, Item
- **LRU algorithm**
- **Auto-populated**



INTEGRATION EXAMPLE: STREAM ANALYTICS

SLIDE 1: STREAM FROM SOURCE TO SINK, WITH OPTIONAL AI



INTEGRATION EXAMPLE: STREAM ANALYTICS

SLIDE 2: SQL "CODE" – SELECT FROM SOURCE INTO SINK

 **cjoakim-sa - Query**
Stream Analytics job

Search (Ctrl+ /) Save Discard Test

Overview Activity log Access control (IAM) Tags Diagnose and solve problems

Inputs (1)
cjoakim-iot-hub

Outputs (1)
cosmosdb-iot-collection

Need help with your query? Check out some of the most common queries.

```
1 SELECT
2 *
3 INTO
4 [cosmosdb-iot-collection]
5 FROM
6 [cjoakim-iot-hub]
```

...

Settings

Locks

Job topology

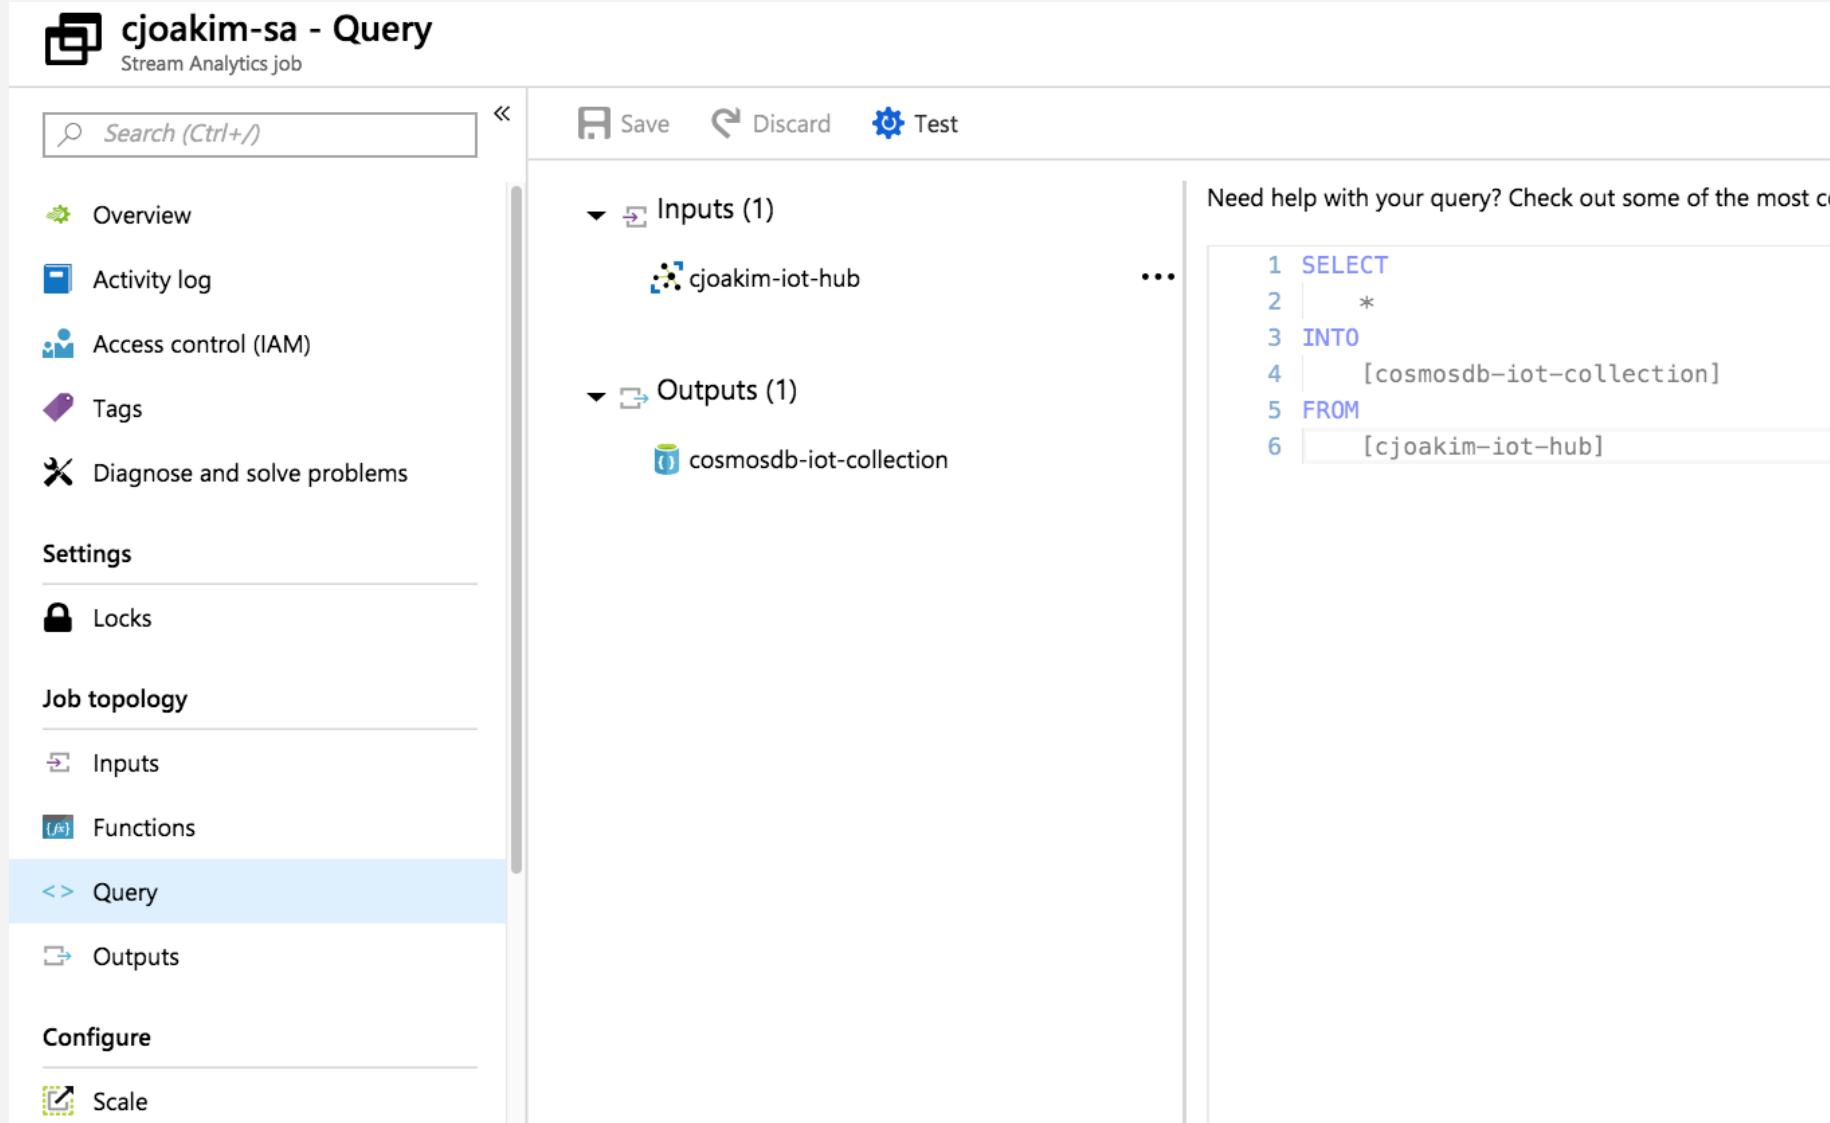
Inputs Functions

Query

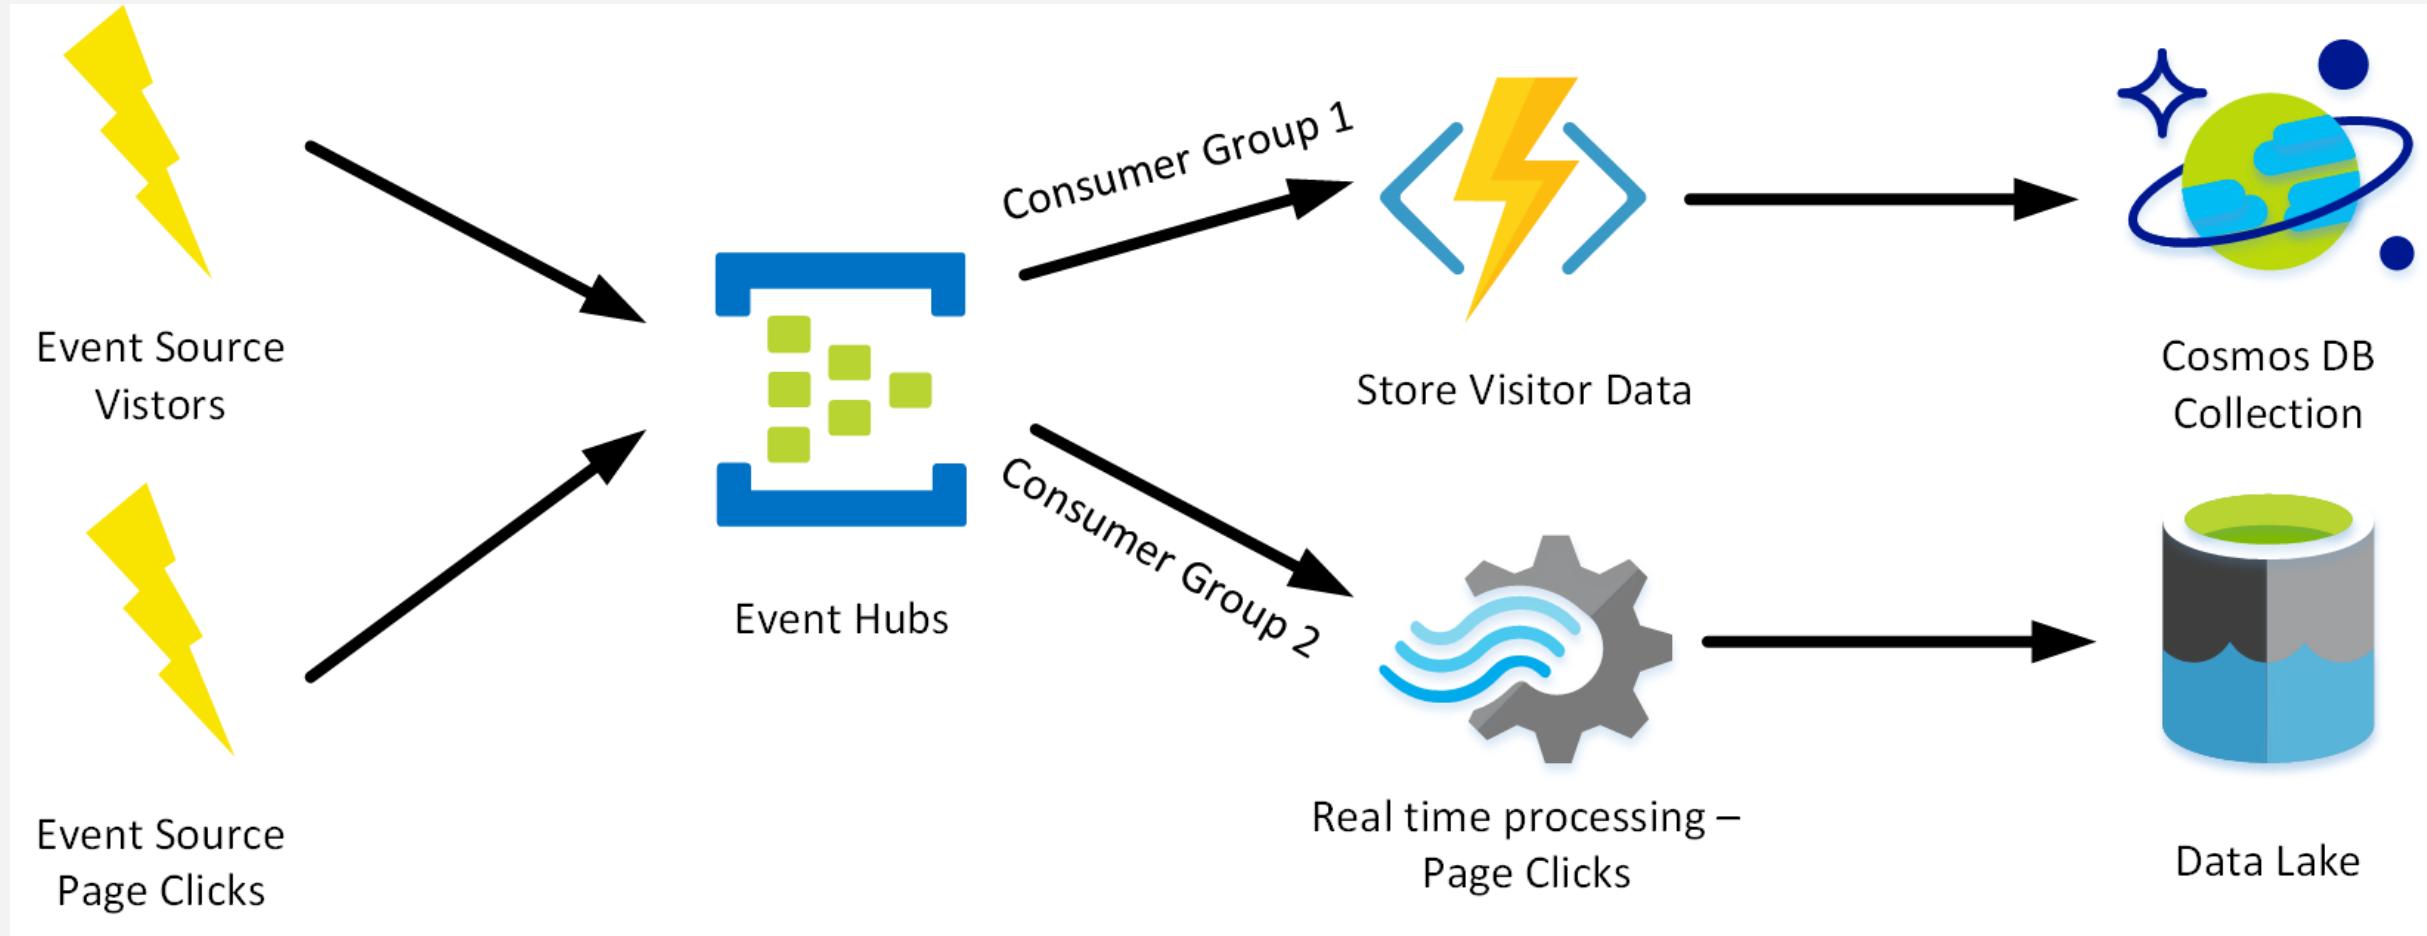
Outputs

Configure

Scale

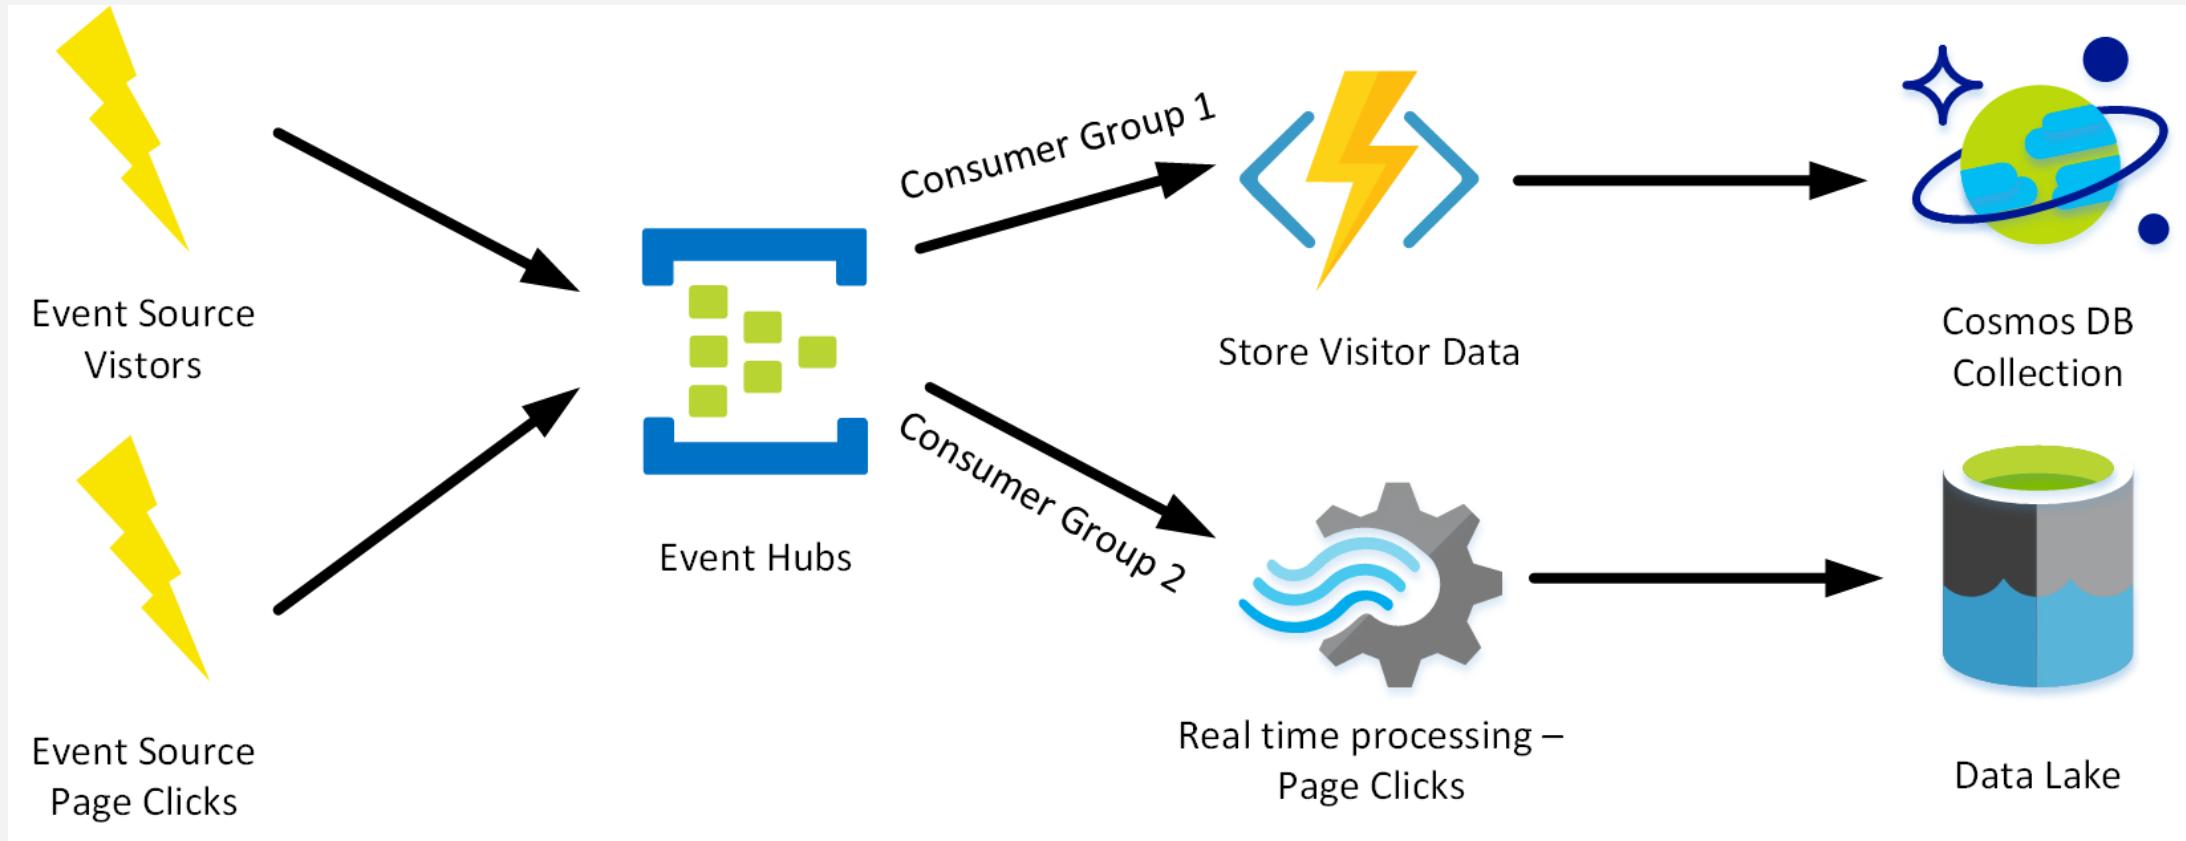


INTEGRATION EXAMPLE: STREAMING



INTEGRATION EXAMPLE: COGNITIVE SEARCH

MULTIPLE POSSIBILITIES FOR SOURCES AND SINKS

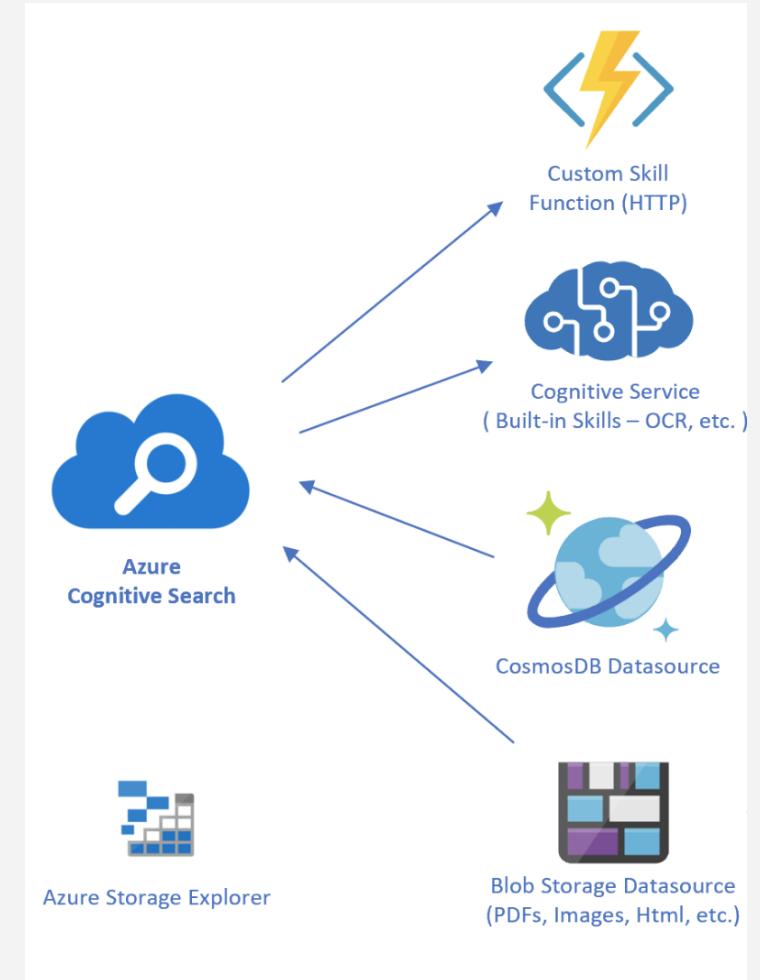


INTEGRATION EXAMPLE: AZURE COGNITIVE SEARCH

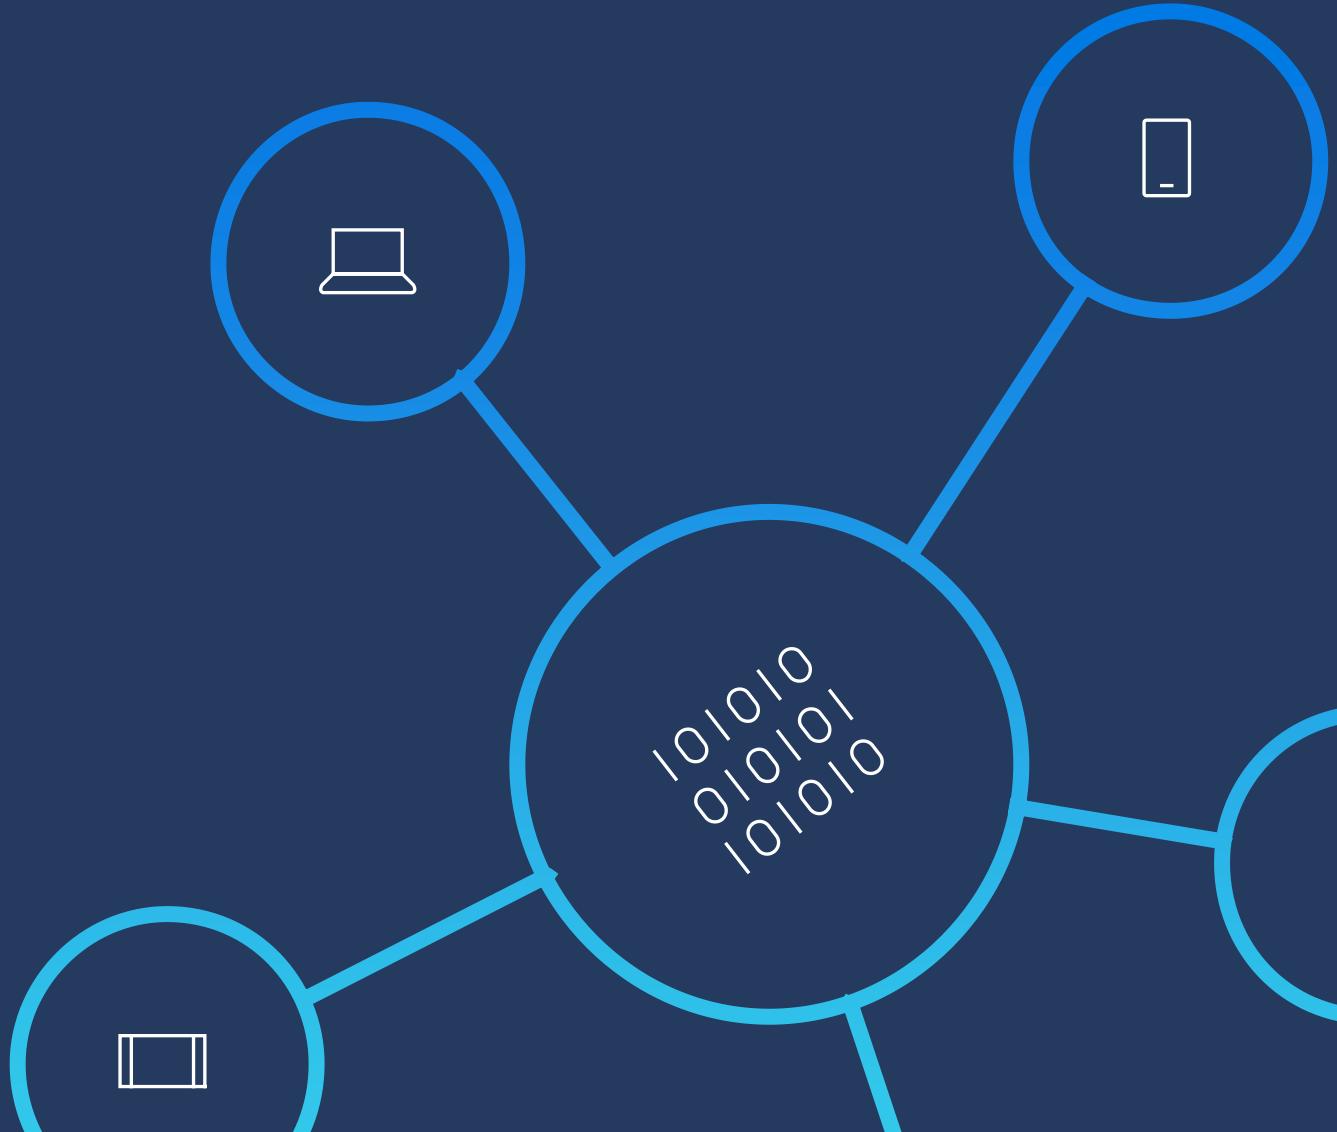
AI-DRIVEN INTELLIGENT SEARCHING



```
"imageDescription": [  
    {"tags": ["person", "road", "outdoor", "sport", "street", "woman", "people", "jumping", "young", "standing", "riding", "city", "playing", "player", "group"], "captions": [{"text": "Shalane Flanagan et al. walking down the street", "confidence": 0.7455881694062344}]}]  
, "imageText": [  
    "B TATA CONSULTANCY SERVICES TATA TCS NEW FLANAGAN 2017 % WOR YORK CITY  
    airbnb"]]
```

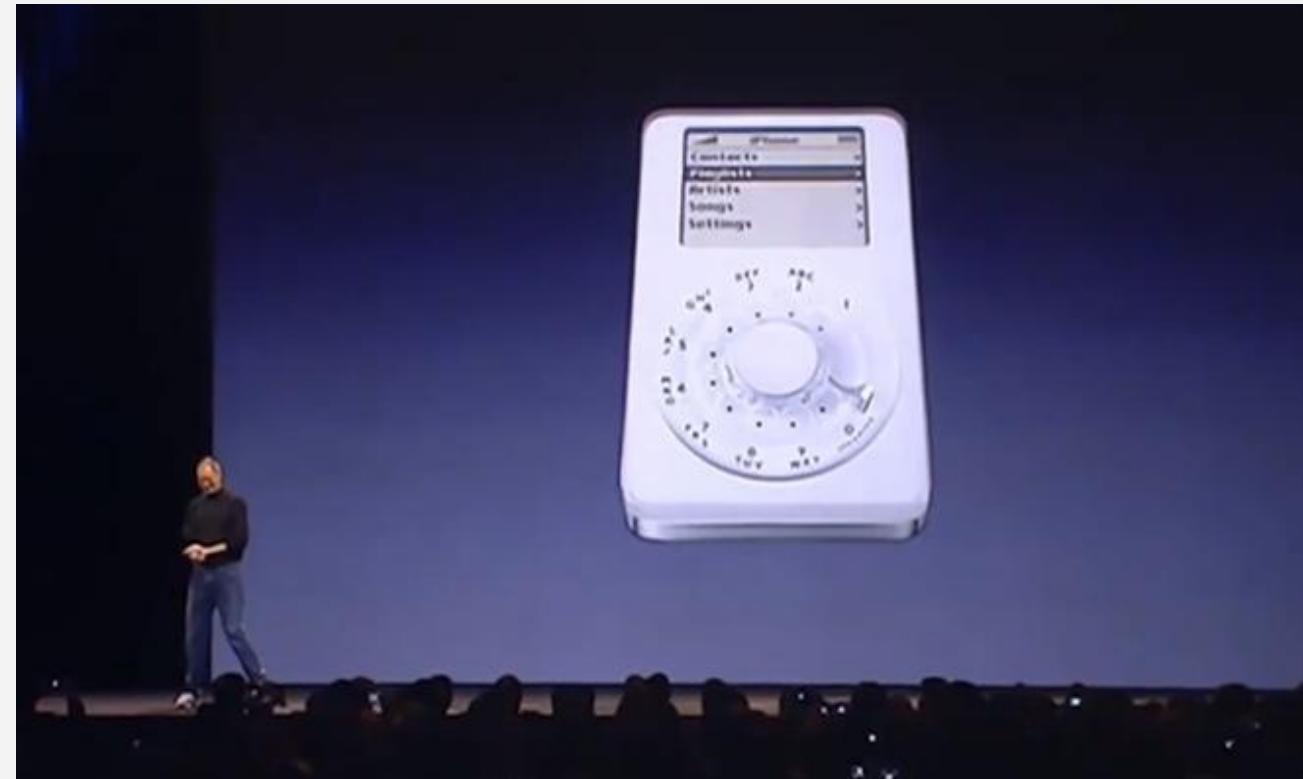


Designing and Developing with CosmosDB



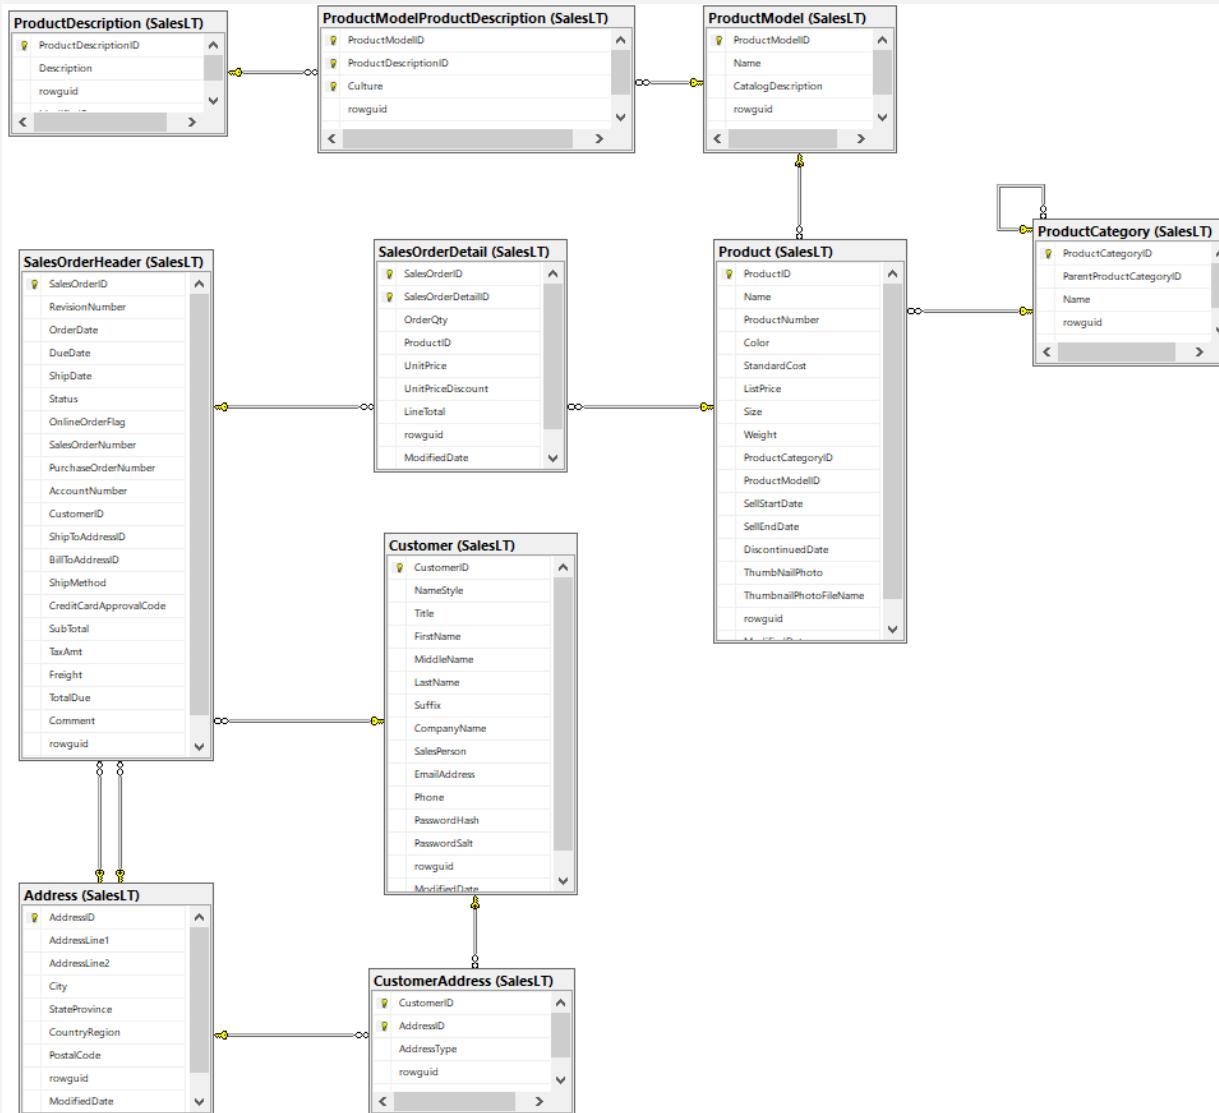
HOW DO I DESIGN THIS NEW THING?

BEWARE OF "DESIGN INERTIA" FROM YOUR PREVIOUS DBMS



DESIGN: ANTI-PATTERNS

COSMOSDB IS NOT RELATIONAL - NO CROSS-CONTAINER JOINS OR RI

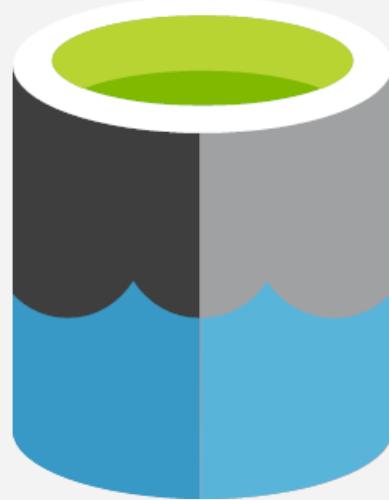


DESIGN: ANTI-PATTERNS

COSMOSDB IS NOT A DATA LAKE; IT'S OLTP NOT OLAP



!=



ONE COMMON AND SUCCESSFUL DESIGN

SINGLE CONTAINER, SMALL DOCUMENTS, PK JOINS

```
{  
  "pk": "XK1123",  
  "doctype": "order",  
  "orderNumber": "XK1123",  
  "customer": "cjoakim",  
  ... other order attributes ...  
}
```

```
{  
  "pk": "XK1123",  
  "doctype": "lineitem",  
  "orderNumber": "XK1123",  
  "number": 1,  
  "item": {"sku": 123, "cost": 12.44, "quantity": 6, "shipping_method": "usps"}  
}
```

```
{  
  "pk": "XK1123",  
  "doctype": "lineitem",  
  "orderNumber": "XK1123",  
  "number": 2,  
  "item": {"sku": 456, "cost": 349.99, "quantity": 1, "shipping_method": "fedex"}  
}
```

SQL Queries for all or parts of the Order

```
select * from c where c.pk = "XK1123"  
select * from c where c.pk = "XK1123" and cdoctype = "order"  
select * from c where c.pk = "XK1123" and cdoctype in ("order", "lineitem")
```

DESIGN BEST PRACTICES, PAGE 1 OF 3

- Use a **Query-Driven or Operation-driven methodology**, not Structure-Driven
- **CQRS** is helpful to first identify your operations and their relative counts
- **Partition Key design is critical** for perf and costs, especially for large containers
 - High-Cardinality and Well-Distributed are desirable
 - But, also a value that you typically query on (i.e. – customerID)
 - Be aware of 20GB limit per Logical Partition
 - Be aware of the 10K RU performance limit per physical partition
 - Be aware of “Hot Partitions”, at time of ingestion, with same Partition Key
 - Consider “Partition Key Joins” (previous slide)
 - Heirarchical (3 part) Partition Keys are currently in preview mode
- **Prefer Small Documents**. The 2MB max is still huge. Consider update costs.
- Beware of Unbounded Embedded data
 - it can lead to large documents and expensive reads and updates

DESIGN BEST PRACTICES, PAGE 2 OF 3

- **It's schemaless.** There is no need to have many "Entity Containers"
 - Use a "**doctype**" or similar attribute to differentiate the documents
- Iterate your design, and **be aware of your RU costs for each operation**
 - Azure Monitor is invaluable for this RU cost analysis
- Leverage Other PaaS services – Synapse Link, Cognitive Search, etc.
- Leverage the **Change Feed for Event-Driven processing**
- Leverage **TTL** to reduce storage and physical partitions
 - Take advantage of "free deletes"
 - **Your minimum RU level can be driven by the size of your container**
 - Each 50GB Physical Partition ~ = 10,000 RU
- Enable Synapse Link to offload data and batch processing to Synapse

DESIGN BEST PRACTICES, PAGE 3 OF 3

- Consider creating “**Materialized Views**”
 - Optimized data structures for fast reads and serving to web layer
 - Create these with Change Feed and Azure Functions, or Synapse Link
- Optimize performance with **Indexing and Composite Indexes**
- **Let Microsoft Help You** with Design and Architecture
 - We want you to be successful, and happy with Azure and CosmosDB

COSMOSDB PROGRAMMING OVERVIEW

- CosmosDB open-source APIs
 - use the open-source SDKs (Mongo/pymongo, Cassandra/datastax, etc.)
- **CosmosDB SQL API**
 - Microsoft SDKs for the major programming languages
 - Java, C#, Python, Node.js, Go, REST
 - Spring Data
 - <https://docs.microsoft.com/en-us/azure/developer/java/spring-framework/how-to-guides-spring-data-cosmosdb>
- Spark
 - <https://docs.microsoft.com/en-us/azure/cosmos-db/sql/create-sql-api-spark?tabs=python>
- Kafka
 - <https://docs.microsoft.com/en-us/azure/cosmos-db/sql/kafka-connector-sink>
- Emulator – For Windows, or Docker Container
 - <https://docs.microsoft.com/en-us/azure/cosmos-db/local-emulator>
- Azure Data Factory
 - <https://docs.microsoft.com/en-us/azure/data-factory/connector-azure-cosmos-db?tabs=data-factory>

COSMOSDB PROGRAMMING – JAVA SDK V4

- <https://docs.microsoft.com/en-us/azure/cosmos-db/sql/sql-api-sdk-java-v4>
- JDK 8+
- Synchronous Client
- Asynchronous Client - Project Reactor, Mono, Flux
- Bulk Executor
 - <https://docs.microsoft.com/en-us/azure/cosmos-db/sql/bulk-executor-java>
- Excellent Insights into Performance
 - Request Unit Costs, Index Utilization, Server time, etc.
- Interesting Features
 - Preferred Regions List
 - Automatic Configurable Retries
 - Partial Document Updates - container.patchItem
- <https://github.com/cjoakim/azure-cosmosdb-java> My samples, Ask for help!

COSMOSDB PROGRAMMING – JAVA SDK V4

Obtain a CosmosAsyncClient

```
if (client == null) {
    client = new CosmosClientBuilder()
        .endpoint(uri)
        .key(key)
        .preferredRegions(AppConfiguration.getInstance().getPreferredRegions())
        .consistencyLevel(ConsistencyLevel.SESSION)
        .contentResponseOnWriteEnabled(true)
        .buildAsyncClient();
    log.warn("client: " + client);

    database = client.getDatabase(this.currentDbName);
    log.warn("client connected to database Id: " + database.getId());
}
return client;
```

COSMOSDB PROGRAMMING – JAVA SDK V4

- Asynch Query Example
- Obtain a FeedResponse Iterable
 - Flux tolterable()
- Iterate the Pages
- Iterate the Page Items
- Aggregate the RU charges

```
public TelemetryQueryResults getAllTelemetry() {  
  
    String sql = "select * from c offset 0 limit 10000";  
    int    pageSize = 100;  
    String continuationToken = null;  
    CosmosQueryRequestOptions queryOptions = new CosmosQueryRequestOptions();  
    TelemetryQueryResults resultsStruct = new TelemetryQueryResults(sql);  
    resultsStruct.start();  
  
    // Execute the SQL query and iterate the paginated result set,  
    // collecting the documents and total RU charge.  
    do {  
        Iterable<FeedResponse<TelemetryEvent>> feedResponseIterator =  
            container.queryItems(sql, queryOptions, TelemetryEvent.class).byPage(  
                continuationToken, pageSize).toIterable(); // Asynch Flux to Iterable  
  
        for (FeedResponse<TelemetryEvent> page : feedResponseIterator) {  
            for (TelemetryEvent doc : page.getResults()) {  
                resultsStruct.addDocument(doc);  
            }  
            resultsStruct.addRequestCharge(page.getRequestCharge());  
            resultsStruct.incrementPageCount();  
            continuationToken = page.getContinuationToken();  
        }  
    }  
    while (continuationToken != null);  
    resultsStruct.stop();  
    return resultsStruct;  
}
```

JAVA: SPRING AND SPRING DATA

A REPOSITORY CLASS

```
16 @Component
17 @Repository
18 public interface TelemetryRepository extends
19     CosmosRepository<TelemetryEvent, String>,
20     TelemetryRepositoryExtensions {
21
22     1 usage
23     Iterable<TelemetryEvent> findByStateCode(String stateCode);
24
25     1 usage
26     Iterable<TelemetryEvent> findByStateCodeAndCountyCode(String stateCode, String countyCode);
27
28     3 usages
29     @Query("select value count(1) from c")
30     long countAllDocuments();
31
32     1 usage
33     @Query( "select c.id, c.pk, c.stateCode, c.countyCode, c.siteNum, c.observationCount, c.nullObservations " +
34         "from c " +
35         "where c.observationCount >= @observationCount " +
36         "and c.nullObservations >= @nullObservations")
37     Iterable<TelemetryEvent> findByObservationCount(
38         @Param("observationCount") long observationCount,
39         @Param("nullObservations") long nullObservations);
40
41 }
```

JAVA: SPRING AND SPRING DATA

A REPOSITORY EXTENSION INTERFACE TO EXTEND THE SPRING DATA API

```
7  /**
8   * This interface was created to extend the TelemetryRepository, which in turn extends
9   * CosmosRepository<TelemetryEvent, String> from the CosmosDB Spring Data SDK.
10  *
11  * This demonstrates how to leverage more of the power of the CosmosDB SQL syntax, by using
12  * "Criteria" objects and an Autowired "CosmosTemplate" object.
13  *
14  * See class TelemetryRepositoryExtensionsImpl in this package, which implements this interface.
15  *
16  * Chris Joakim, Microsoft, September 2022
17  */
18 public interface TelemetryRepositoryExtensions {
19
20     public Iterable<TelemetryEvent> findByStateCodeAndSiteNumbers(String stateCode, ArrayList<String> siteNumbers);
21 }
```

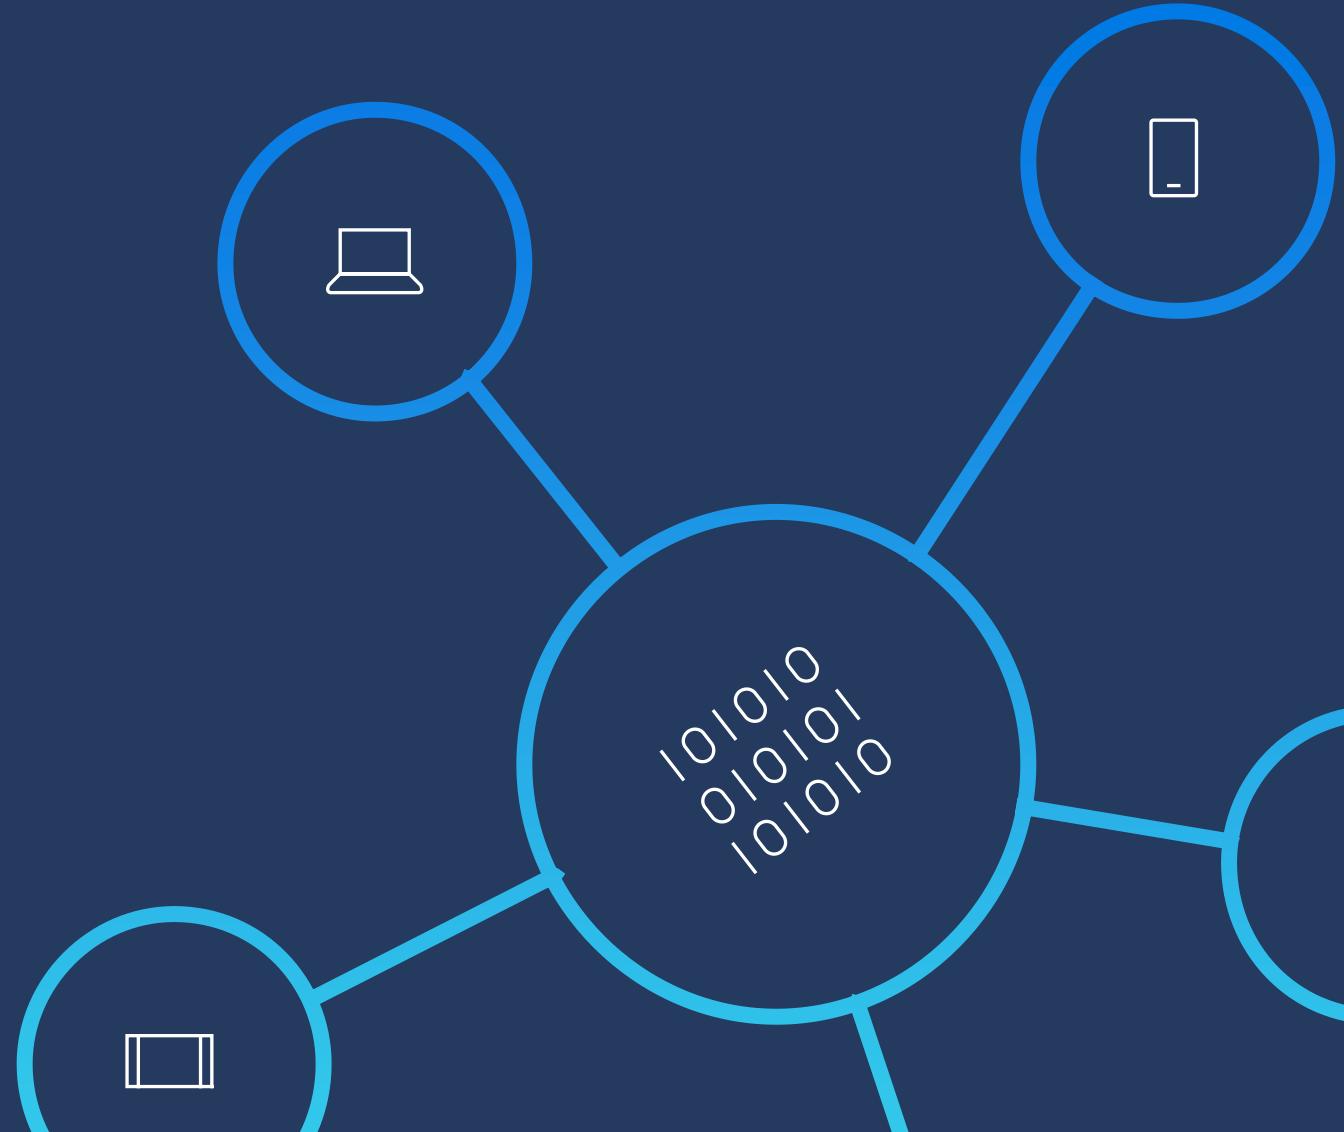
JAVA: SPRING AND SPRING DATA

A REPOSITORY EXTENSION INTERFACE IMPLEMENTATION

- @Autowired
CosmosTemplate
- Criteria Objects
- Criteria.IS_EQUAL
- Criteria.IN
- Criteria.AND
- CosmosQuery
- template.find()

```
24  @Slf4j
25  public class TelemetryRepositoryExtensionsImpl implements TelemetryRepositoryExtensions {
26      3 usages
27      private CosmosTemplate template;
28
29      @Autowired
30      public TelemetryRepositoryExtensionsImpl(CosmosTemplate t) {
31          super();
32          this.template = t;
33          log.warn("TelemetryRepositoryExtensionsImpl constructor, template: " + this.template);
34      }
35  ①↑
36
37      1 usage
38      public Iterable<TelemetryEvent> findByStateCodeAndSiteNumbers(String stateCode, ArrayList<String> siteNumbers) {
39
40          String containerName = AppConstants.TELEMETRY_CONTAINER_NAME;
41
42          Criteria criteria1 = Criteria.getInstance(
43              CriteriaType.IS_EQUAL, subject: "stateCode", Collections.singletonList(stateCode),
44              Part.IgnoreCaseType.NEVER);
45          Criteria criteria2 = Criteria.getInstance(
46              CriteriaType.IN, subject: "siteNum", Collections.singletonList(siteNumbers),
47              Part.IgnoreCaseType.NEVER);
48          Criteria allCriteria = Criteria.getInstance(CriteriaType.AND, criteria1, criteria2);
49
50          CosmosQuery query = new CosmosQuery(allCriteria);
51          return template.find(query, TelemetryEvent.class, containerName);
52      }
53  }
```

Questions & Discussion





Thank You!